Total points 6



Which three of the following are features of the Apache Spark application?

1/1 point

Distributed execution engine

Correct

The Apache Spark core engine is a distributed execution engine.

In-memory processing

Correct

Apache Spark is a parallel processing framework that supports in-memory processing to boost the performance of big-data analytics applications.

Disk-based processing

Parallel Processing Framework

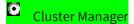
Correct

Apache Spark is a parallel processing framework that supports in-memory processing to boost the performance of big-data analytics applications.

2. Question 2

Which one of the following objects is responsible for allocating resources across applications in an Apache Spark pool?

1/1 point



Nodes

SparkContext

C Executors

Correct

The SparkContext can connect to the cluster manager, which allocates resources across applications.

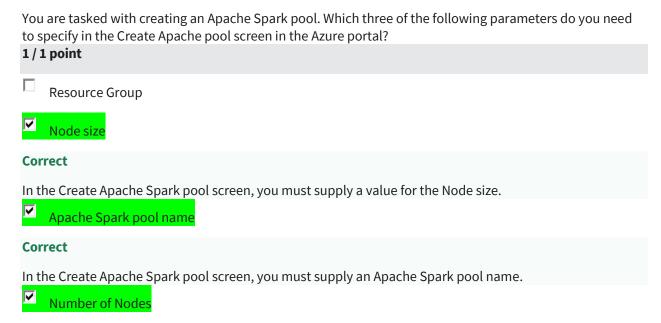
3. Question 3

Apache Spark pools in Azure Synapse Analytics are compatible with which of the following types of storage?

Azure Storage
Correct
Spark pools in Azure Synapse Analytics are compatible with Azure Storage.
SQL Storage
Azure Data Lake Generation 2 Storage
Correct
Spark pools in Azure Synapse Analytics are compatible with Azure Data Lake Generation 2 Storage.
Azure Data Lake Generation 1 Storage
4. Question 4You need to manage an end-to-end big data project using one single platform. Which of the following data services is best suited to this task?1/1 point
C Apache Spark for Azure Synapse
Azure HDInsight
Azure Databricks Apache Spark
Correct
Azure Databricks provides the ability to create and manage an end-to-end big data/data science project using one platform.
5. Question 5Which of the following is an element of an Apache Spark Pool in Azure Synapse Analytics?1/1 point
C Apache Spark Console
C Azure HDInsight
Spark Instance
Correct
The definition of an Apache Spark pool is that, when instantiated, it is used to create an Apache Spark

The definition of an Apache Spark pool is that, when instantiated, it is used to create an Apache Spark instance that processes data.

6. Question 6



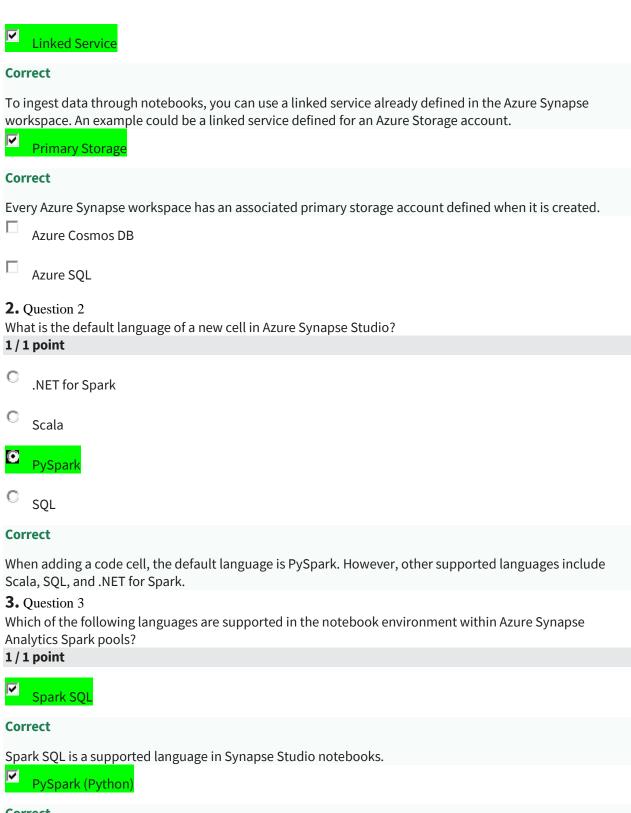
In the Create Apache Spark pool screen, you must specify the number of Nodes to use in the Pool.

Knowledge check

Total points 6

1. Question 1

You need to ingest data through Apache Spark notebooks. Which two of the following features can you use to carry out this task?



PySpark is a supported language in Synapse Studio notebooks.

Spark (Scala)

Scala is a supported language in Synapse Studio notebooks. NET Spark (CB) Correct Net Spark is a supported language in Synapse Studio notebooks. JSON YAML 4. Question 4 Your Azure Studio notebook needs to be able to reference data or variables directly using different languages. Which of the following actions do you need to perform to enable this? 1/1 point Use a magic command for that language. You don't need to do anything as you can reference data or variables directly using different languages in an Azure Synapse Studio notebook. Create a temporary table so that it can be referenced across different languages. Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark Apache Spark
.Net Spark is a supported language in Synapse Studio notebooks. JSON YAML 4. Question 4 Your Azure Studio notebook needs to be able to reference data or variables directly using different languages. Which of the following actions do you need to perform to enable this? 1/1 point Use a magic command for that language. You don't need to do anything as you can reference data or variables directly using different languages in an Azure Synapse Studio notebook. Create a temporary table so that it can be referenced across different languages. Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
JSON YAML 4. Question 4 Your Azure Studio notebook needs to be able to reference data or variables directly using different languages. Which of the following actions do you need to perform to enable this? 1/1 point Use a magic command for that language. You don't need to do anything as you can reference data or variables directly using different languages in an Azure Synapse Studio notebook. Create a temporary table so that it can be referenced across different languages. Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
4. Question 4 Your Azure Studio notebook needs to be able to reference data or variables directly using different languages. Which of the following actions do you need to perform to enable this? 1/1 point Use a magic command for that language. You don't need to do anything as you can reference data or variables directly using different languages in an Azure Synapse Studio notebook. Create a temporary table so that it can be referenced across different languages. Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
Your Azure Studio notebook needs to be able to reference data or variables directly using different languages. Which of the following actions do you need to perform to enable this? 1/1 point Use a magic command for that language. You don't need to do anything as you can reference data or variables directly using different languages in an Azure Synapse Studio notebook. Create a temporary table so that it can be referenced across different languages. Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
You don't need to do anything as you can reference data or variables directly using different languages in an Azure Synapse Studio notebook. Create a temporary table so that it can be referenced across different languages. Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
in an Azure Synapse Studio notebook. Create a temporary table so that it can be referenced across different languages. Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
Create a new Notebook. Correct You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
notebook. If you wish to do this using Spark, you should first create a temporary table so that it can be referenced across different languages. 5. Question 5 Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
Azure Synapse Studio notebooks are based on which one of the following? 1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
1/1 point Dedicated SQL pool Apache Spark Apache Spark pool
Apache Spark Apache Spark pool
C Apache Spark pool
Correct
Correct! Azure Synapse Studio notebook is purely Apache Spark based.
6. Question 6 Which one of the following actions should you take to save all notebooks in Azure Synapse studio? 1/1 point

9	Select the Publish all button on the workspace command bar.
О	Select the Publish button on the notebook command bar.
С	Press CTRL + S.
Cor	rect
Tos	ave all notebooks in your workspace, select the Publish all button on the workspace command bar.

Total points 6

1. Question 1

Which three of the following are features of DataFrames?

1/1 point

DataFrames are a collection of data organized into named Rows.

DataFrames optimize execution plans on queries that will access the data held in the DataFrame.

Correct

DataFrames optimize execution plans on gueries that will access the data held in the DataFrame.

DataFrames are a collection of data organized into named columns.

Correct

DataFrames are a collection of data organized into named columns.

DataFrames enable Apache Spark to understand the schema of the data.

Correct

DataFrames enable Apache Spark to understand the schema of the data.

2. Question 2

You input the following Python code snippet into your workspace:

new_rows = [('CA',22, 45000),("WA",35,65000),("WA",50,85000)]

demo_df = spark.createDataFrame(new_rows, ['state', 'age', 'salary'])

demo_df.show()

The variable named *demo_df* above is used to do which of the following?

1/1 point

- It uses the *spark.createDataFrame* method referencing the *new_rows* variable in the first parameter. The second parameter defines the column heading names for the DataFrame as state, age, and salary.
- It uses the *spark.createDataFrame* method to create a variable named new_rows which will store the values '*state*', '*age*', and '*salary*'.
- It uses the *spark.createDataFrame* method and creates a variable named *new_rows* which creates the data in the code segment to store in the DataFrame.

Correct

The second variable named *demo_df* uses the *spark.createDataFrame* method, which refers to the *new_rows* variable in the first parameter. The second parameter defines the column heading names for the DataFrame as *state*, *age*, and *salary*.

3. Question 3

You input the following Python snippet into your code:

from azureml.opendatasets import NycTlcYellow

data = NycTlcYellow()

data_df = data.to_spark_dataframe()

display(data_df.limit(10))

What is the purpose of the <code>display(data_df.limit(10))</code> method?

1/1 point

- Return 10 rows of data from the data_df variable
- Return batches of 10 rows of data from the *data_df* variable until all records are returned.
- Limit the Dataframe to retrieving 10 rows of data from the *NycTLcYellow* data source.

Correct

10 rows of data are returned from the *data_df* variable using the display method.

4. Question 4

Select the correct series of steps to flatten nested structures and explode arrays with Apache Spark

1/1 point

()

- 1. Define a function
- 2. Flatten nested schema
- Explode Arrays
- 1. Flatten child nested schema

O	
	1. Define a function
	2. Flatten nested schema
	3. Flatten child nested schema4. Explode Arrays
O	4. Explode Arrays
	1. Define a function
	2. Flatten child nested schema3. Flatten nested schema
	4. Explode Arrays
	rrect
Thi	s is the correct sequence of steps to flatten nested structures and explode arrays with Apache Spark.
5. (Question 5
	ich of these actions should you perform to flatten a nested schema?
	333333333333 / 1 point
	Explode Arrays.
	Load a CSV file.
V	Create a Parquet file.
	s should not be selected
Cre	eating a Parquet file is not required to flatten a nested schema. It is a file format.
6.	Question 6
	ich two of the following actions do DataFrames perform?
	1 point
	Process data only in streaming data architecture.
	Process data only in batch data architecture.
	Extract large volumes of data from an SQL Database only.
V	Process data in either batch or streaming data architecture.
	rrect
	taFrames can process data in either batch or streaming data architecture.
V	Extract large volumes of data from a wide variety of data sources.

DataFrames are designed to extract large volumes of data from a wide variety of data sources.

Total points 6

1. Question 1

The interoperability between Apache Spark and SQL helps you to directly explore and analyze which three of the following types of files?

1/1 point



Correct

The interoperability between Apache Spark and SQL helps you to directly explore and analyze CSV files.



Correct

The interoperability between Apache Spark and SQL helps you to directly explore and analyze TSV files.

L YAML

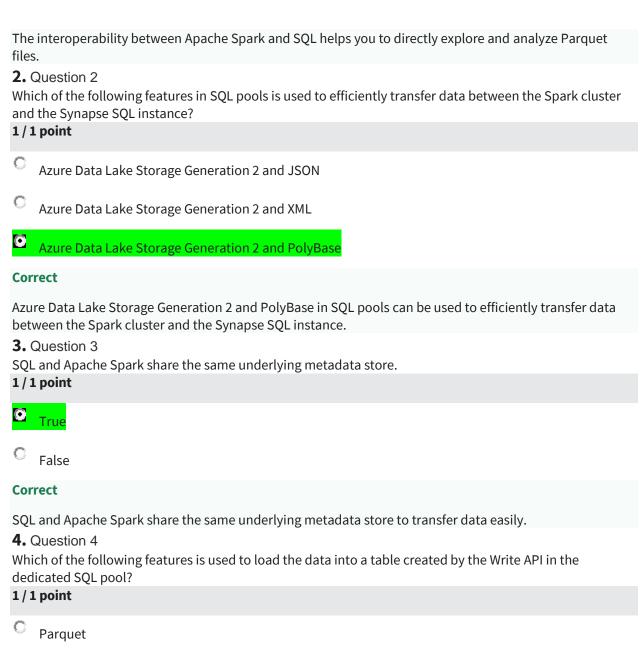


Correct

The interoperability between Apache Spark and SQL helps you to directly explore and analyze JSON files.



Correct



ORC

Son



Correct

The Write API creates a table in a dedicated SQL pool. It then invokes PolyBase to load the data into the table that was created.

5. Question 5

You have a requirement to transfer data to a dedicated SQL pool that is outside of the Azure Synapse Analytics workspace. Which form of Authentication can be used to complete this task?		
0 / 1 point		
SQL Authentication Only		
C Azure AD and SQL Authentication		
C None of the above		
C Azure AD only		
Incorrect		
To establish and transfer data to a dedicated SQL pool that is outside of the workspace without Azure AD, you can use SQL Authentication if AD is not available		
6. Question 6		
The Azure Synapse Apache Spark to Synapse SQL connector supports which one of the following languages?		
1/1 point		
C .Net		
C SQL		
Scala		
C Python		
Correct		
The connector uses Scala to integrate Apache Spark pools with dedicated SQL pools in Azure Synapse Analytics.		

Total points 6

4	\sim		- 4
	()	uestion	1

1. Question i
The Apache Spark history server can be accessed directly from which of the following Synapse Studio Tabs.
0.75 / 1 point
Develop
Manage
Monitor Monitor
Correct

If you want to navigate to the Apache Spark History server, you can navigate to the Azure Synapse Analytics Studio environment and go to the Monitor tab. In the Monitor tab, you can select Apache Spark Applications and from there select Spark history server.



This should not be selected

To open the Apache Spark History server you can navigate to the Data tab, where if you create a notebook and read a DataFrame you can go to the bottom of the page and find the Spark History Server known as the Spark UI.

2. Question 2

Which of the following features can benefit from query optimization through Catalyst when optimizing Apache Spark jobs in Azure Synapse Analytics?

1/1 point

C Resilient Distributed Datasets (RDDs)



O Notebooks

Correct

DataFrames provide query optimization through Catalyst.

Question 3

You need to specify the minimum number of nodes for an Apache Spark pool with autoscaling. What is the minimum number of nodes allowed?

1/1 point

0 :

0

2

C 4
Correct
The minimum number of nodes allowed is 3.
4. Question 4 You scale down Apache Spark pools in Azure Synapse Analytics. What happens to existing nodes?
1/1 point
Nodes to be scaled down will be put in a decommissioned state.
Correct
When you scale down Apache spark pools, Nodes to be scaled down will be put in a decommissioned state. Pending jobs will be in a waiting state and scheduled for execution on fewer nodes.
Correct
When you scale down Apache spark pools, pending jobs will be in a waiting state and scheduled for execution on fewer nodes.
Jobs that are still running will continue to run until completion.
Correct
When you scale down Apache spark pools, jobs that are still running will continue to run until completion.
Nodes to be scaled down will be shut down immediately regardless of current state.
Pending jobs will be lost.
5. Question 5
Which of the following tasks can you perform to optimize an Apache Spark job? 0.66666666666666666666666666666666666
Remove the Apache Spark Pool.
This should not be selected
This is not a valid option as removing the Apache Spark pool removes the capability to process Apache Spark jobs.
Remove all nodes.
Correct
This is not a valid option as removing all nodes would mean there is no compute power to process the Apache Spark jobs.
Use bucketing.

6. Question 6

You encounter a slow performing shuffle job. Which of the following is a possible cause?

1/1 point

- C Bucketing
- C Enablement of autoscaling



Correct

Correct! The slow performance is due to asymmetry in your job data.

Test prep

Latest Submission Grade 100%

1. Question 1

Apache Spark pools in Azure Synapse Analytics benefit from which four of the following?

1/1 point

Pre-Loaded Anaconda libraries

Correct

Over 200 Anaconda libraries are pre-installed on the Spark pool in Azure Synapse Analytics.

REST APIS

Correct

In order to monitor and submit jobs remotely, you can use Apache Livy as a Rest API for the Spark job server.

Apache Livy is a service that enables interaction with an Spark cluster over a REST interface to enable:

- submission of Spark jobs
- snippets of Spark code
- synchronous or asynchronous result retrieval
- and Spark Context management

Autoscale

Correct

Provide support for autoscale, so that pools can be scaled up or down as required, by adding or removing nodes.

Real-time co-authoring (both authors see the changes in real-time)

Support for third party IDEs

Correct

Azure Synapse Analytics provides an IDE for IntelliJ to create and submit applications to the Apache Spark pool.

2. Question 2

Which one of the following resources should you create first when building an Apache Spark pool in Azure Synapse Analytics?

1/1 point

C A notebook

A SQL Database

A workspace

Correct

An Azure Synapse Analytics workspace must be built before creating a new Apache Spark pool.

3. Question 3

Which one of the following resources requires a unique Spark Pool name when creating an Apache Spark pool in Azure Synapse Analytics?

The workspace
The resource group
C Azure
The subscription Correct
An Apache Spark pool name must adhere to the following specifications:
 Contain letters or numbers only, Contain 15 or less characters, Start with a letter, Contain no reserved words, and be unique in the workspace.
4. Question 4 You need to create an embedded Spark capability that can reside alongside, and integrate with, other Azure data warehouse and integration services. Which one of the following solutions can provide these functionalities?
1/1 point
Azure HDInsight Apache Spark for Azure Synapse
Azure Databricks
Apache Spark Correct
Apache Spark is an embedded Apache Spark capability within Azure Synapse Analytics residing on the same platform that contains data warehouses and data integration capabilities, as well as integrating with other Azure services.
5. Question 5 You can use the magic command %%spark to execute which one of the following types of query against a spark context?
1/1 point
SparkSQL
Python
.Net for C#
Scala Scala

The magic command %%spark runs a Scala query against spark context.

6. Question 6

You are using IDE-style IntelliSense to assist with your code completion. Which two of the following languages does the IntelliSense support for Syntax Code Completion?

□ .Net for Spark (C#) Spark (Scala)
Correct The IntelliSense supports Syntax Code Completion for Spark (Scala).
SparkSQL PySpark (Python) Correct
The IntelliSense supports Syntax Code Completion for PySpark (Python).
7. Question 7 You run a notebook in Apache Spark pool for the first time. How many minutes will it take for Azure Synapse Analytics to create a new session? 1/1 point
10 - 15 Minutes 3 - 5 Minutes 30 - 60 Seconds
Correct The first time you run a notebook in a Spark pool, Synapse creates a new session. This can take approximately 3-5 minutes.
8. Question 8 What is a DataFrame?
1/1 point A parquet file
A csv file A data structure
Correct A DataFrame creates a data structure and it's one of the core data structures in Apache Spark.
9. Question 9 You need to load data from different file types into an Apache Spark DataFrame. Which four of the following Azure services can you load the data from?
1 / 1 point Primary Storage Account
You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.
Azure Storage Account
Correct You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.
Azure Data Lake Store Generation 2

You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.

Serverless SQL Pool

Dedicated SQL pool

Correct

You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.

10. Question 10

You need to optimize the Apache Spark Jobs in Azure Synapse Analytics. However, first you need to consider the cluster configuration for the workload you're running on that cluster. Which one of the following resources can benefit from query optimization through Catalyst?

1/1 point

Resilient Distributed Datasets (RDDs)



Notebooks

Correct

DataFrames provide query optimization through Catalyst.

Test prep

Latest Submission Grade 85%

1. Question 1

The Azure Synapse Apache Spark to Synapse SQL connector is designed to efficiently transfer data between which of the following?

1/1 point

- Dedicated Apache Spark pools and Serverless SQL pools in Azure Synapse
- Serverless Apache Spark pools and Serverless SQL pools in Azure Synapse
- Serverless Apache Spark pools and Dedicated SQL pools in Azure Synapse

Correct

The Azure Synapse Apache Spark to Synapse SQL connector is designed to efficiently transfer data between serverless Apache Spark pools and dedicated SQL pools in Azure Synapse.

2. Question 2

You need to develop and execute transformation pipelines within the Azure Synapse Studio notebook experience. Which three of the following languages will help you in your task?

1/1 point

 \square JSON



Correct

You can attach a SQL or Apache Spark pool, and develop and execute transformation pipelines using Spark (Scala).



Correct

You can attach a SQL or Apache Spark pool, and develop and execute transformation pipelines using SparkSQL.



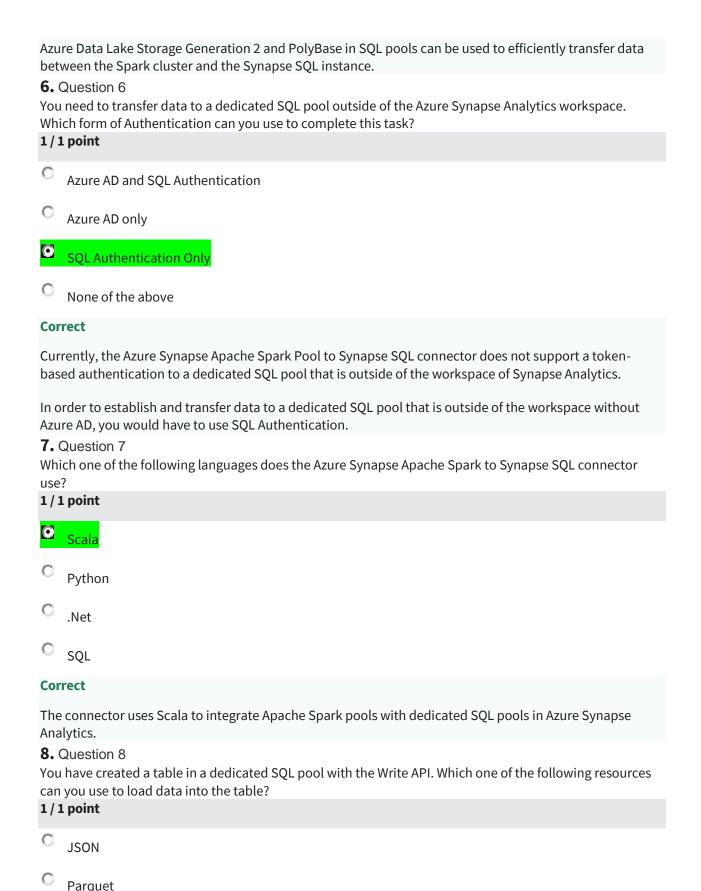
Correct

You can attach a SQL or Apache Spark pool, and develop and execute transformation pipelines using Python.

3. Question 3

Which two of the following prerequisites are necessary to successfully authenticate an account between two different systems?

The account used needs to be a member of the Storage Blob Data Contributor role on the default storage account.
Correct
The account used needs to be a member of the Storage Blob Data Contributor role on the default storage account.
The account used needs to be a member of Storage Blob Data Contributor role in the database or SQL pool from which you to transfer data to or from.
The account used needs to be a member of db_exporter role in the database or SQL pool from which you to transfer data to or from.
Correct
The account used needs to be a member of db_exporter role in the database or SQL pool from which you to transfer data to or from.
The account used needs to be a member of the db_exporter role on the default storage account.
4. Question 4In which of the following scenarios are you less likely to require import statements when transferring data between a dedicate SQL and Apache Spark pool?0 / 1 point
• When relying on the use of token-based authentication.
When using the Azure integrated notebook experience.
When using the Apache Spark PySpark connector.
Incorrect
Try going back and reviewing Query pools and manage workloads in Azure Synapse Analytics.
5. Question 5 Which of the following SQL pools can you use to efficiently transfer data between the Apache Spark cluster and the Azure Synapse SQL instance with the Azure Synapse Apache Spark pool to Synapse SQL connector?
1 / 1 point
Azure Data Lake Storage Generation 2 and PolyBase
C Azure Data Lake Storage Generation 2 and JSON
C Azure Data Lake Storage Generation 2 and XML
Correct



C ORC
PolyBase
Correct
The Write API creates a table in a dedicated SQL pool. It then invokes PolyBase to load the data into the table that was created.
9. Question 9 Which two of the following Azure Synapse Studio tabs can you use to access the Apache Spark history server?
0.5 / 1 point
Manage
Develop
This should not be selected
Try going back and reviewing Query pools and manage workloads in Azure Synapse Analytics.
Monitor
Data Data
Correct
To open the Apache Spark History server, you can navigate to the Data tab, where if you create a notebook and read a DataFrame you can go to the bottom of the page and find the Spark History Server known as the Spark UI.
10. Question 10
Which three of the following actions occur within existing nodes when you scale down Apache Spark pools? 1/1 point
Nodes to be scaled down will be put in a decommissioned state.
Correct
Nodes will be placed in a decommissioned state. Pending jobs will be in a waiting state and scheduled for execution on fewer nodes.
Correct When you scale down Apache spark pools, pending jobs will be placed in a waiting state and scheduled for execution on fewer nodes.
Nodes to be scaled down will be shut down immediately regardless of current state.
Pending jobs will be lost.



Jobs that are still running will continue to run and finish.

Correct

Running jobs will continue until completion.

Course practice exam

Latest Submission Grade 93.33%



Can you identify three features of Apache Spark?

1/1 point

In-memory processing

Correct

Apache Spark is a parallel processing framework that supports in-memory processing to boost the performance of big-data analytics applications.

Distributed execution engine

Correct

The Apache Spark core engine is a distributed execution engine.

Parallel Processing Framework

Correct

Apache Spark is a parallel processing framework that supports in-memory processing to boost the performance of big-data analytics applications.

Disk-based processing

2. Question 2

Which object is responsible for allocating resources across Apache Spark applications in Azure Synapse Analytics?

1/1 point

Cluster Manager

SparkContext

Nodes

Executors

Correct

The SparkContext can connect to the cluster manager, which allocates resources across applications.

3. Question 3

Spark pools in Azure Synapse Analytics are compatible with which two of the following storage types?

1/1 point

SQL Storage

Azure Data Lake Generation 2 Storage

Correct

Spark pools in Azure Synapse Analytics are compatible with Azure Data Lake Generation 2 Storage.

Azure Data Lake Generation 1 Storage

Azure Storage

Correct

Spark pools in Azure Synapse Analytics are compatible with Azure Storage.

4. Question 4

Which of the following solutions can you utilize to create and manage an end-to-end big data or data science project using one single platform?
1/1 point
Azure HDInsight
Apache Spark
Azure Databricks
C Apache Spark for Azure Synapse
Correct
Azure Databricks provides the ability to create and manage an end-to-end big data/data science project using one platform.
5. Question 5
What is the default language of a new cell in Azure Synapse Studio? 1/1 point
c .
2Ár
Scala
O .NET for Spark
PySpark
Correct When adding a code cell, the default language is PySpark. However, other supported languages include
Scala, SQL, and .NET for Spark.
6. Question 6
You deploy the magic command %%spark. What type of query will this command execute against Spark
context?
1/1 point
Python
Scala Scala
SparkSQL
.Net for C#
Correct The magic command %%spark runs a Scala query against spark context.
7. Question 7
Which of the following actions do you need to perform in order to directly reference data or variables in Azure Synapse Studio notebook using different languages?
1/1 point
Create a new Notebook.
Do Nothing. You can reference data or variables directly using different languages in an Azure Synapse Studio notebook.
C
Use a magic command for that language.

Create a temporary table so that it can be referenced across different languages.

Correct

You cannot reference data or variables directly using different languages in an Azure Synapse Studio notebook. If you wish to do this using Spark, you must first create a temporary table so that it can be referenced across different languages.

8. Question 8

What are DataFrames?

1/1 point

DataFrames are a collection of data organized into named Rows.

DataFrames optimize execution plans on queries that will access the data held in the DataFrame

DataFrames optimize execution plans on queries that will access the data held in the DataFrame.

DataFrames enable Apache Spark to understand the schema of the data.

Correct

DataFrames enable Apache Spark to understand the schema of the data.

DataFrames are a collection of data organized into named columns.

Correct

DataFrames are a collection of data organized into named columns.

9. Question 9

You input the following Python snippet into your code:

new_rows = [('CA',22, 45000),("WA",35,65000),("WA",50,85000)]

demo_df = spark.createDataFrame(new_rows, ['state', 'age', 'salary'])

demo_df.show()

The variable named *demo_df* performs which of the following actions?

0 / 1 point

It uses the *spark.createDataFrame* method referencing the *new_rows* variable in the first parameter. The second parameter defines the column heading names for the DataFrame as *state*, *age*, and *salary*.

It uses the *spark.createDataFrame* method and creates a variable named *new_rows* which creates the data in the code segment to store in the DataFrame.

It uses the *spark.createDataFrame* method to create a variable named *new_rows* which will store the values *state, age,* and *salary.*

Incorrect

Try going back and reviewing Big Data Engineering.

10. Question 10

You need to flatten nested structures and explode arrays with Apache Spark. What series of steps should you perform to complete these tasks?

1/1 point

 \bigcirc

- 1. Define a function
- 2. Flatten child nested schema
- 3. Flatten nested schema
- 4. Explode Arrays

О

- 1. Define a function
- 2. Flatten nested schema
- 3. Flatten child nested schema
- 4. Explode Arrays

•

- 1. Define a function
- 2. Flatten nested schema
- Explode Arrays
- 4. Flatten child nested schema

Correct

The correct order to flatten nested structures and explode arrays with Apache Spark are:

- 1. Define a function
- 2. Flatten nested schema
- 3. Explode Arrays
- 4. Flatten child nested schema

11. Question 11

You can use the Azure Synapse Apache Spark to Synapse SQL connector to transfer data between which of the following?

1/1 point

Serverless Apache Spark pools and Dedicated SQL pools in Azure Synapse.

Serverless Apache Spark pools and Serverless SQL pools in Azure Synapse.

Dedicated Apache Spark pools and Serverless SQL pools in Azure Synapse.

Correct

The Azure Synapse Apache Spark to Synapse SQL connector is designed to efficiently transfer data between serverless Apache Spark pools and dedicated SQL pools in Azure Synapse.

12. Question 12

Which of the following role memberships are required to successfully authenticate between two systems in Azure Synapse Analytics?

1/1 point

The account used needs to be a member of Storage Blob Data Contributor role in the database or SQL pool from which you to transfer data to or from.

The account used needs to be a member of db_exporter role in the database or SQL pool from which you to transfer data to or from.

Correct

The account used needs to be a member of db_exporter role in the database or SQL pool from which you to transfer data to or from.

The account used needs to be a member of the Storage Blob Data Contributor role on the default storage account.
Correct Try going back and reviewing Query pools and manage workloads in Azure Synapse Analytics.
13. Question 13 You have a requirement to transfer data to a dedicated SQL pool that is outside of the workspace of Synapse Analytics. Which form of Authentication can you use to complete this task? 1/1 point
Azure AD and SQL Authentication.
SQL Authentication Only.
None of the above.
C Azure AD only.
Correct To establish and transfer data to a dedicated SQL pool that is outside of the workspace you can use SQL Authentication or Azure AD if available
14. Question 14 What is the minimum number of nodes allowed when creating an Apache Spark pool with Autoscaling? 1/1 point
1 C 4
Correct The minimum number of nodes allowed is 3.
15. Question 15 What three actions occur within existing nodes in Azure Synapse Analytics when you scale down Apache Spark pools?
1/1 point
Jobs that are still running will continue to run and finish. Correct
In Azure Synapse when you scale down Apache spark pools Jobs that are still running will continue to run and finish.
Pending jobs will be lost.
Nodes to be scaled down will be shut down immediately regardless of current state.
Pending jobs will be in a waiting state and scheduled for execution on fewer nodes.
Correct In Azure Synapse when you scale down Apache spark pools pending jobs will be in a waiting state and scheduled for execution on fewer nodes.
Nodes to be scaled down will be put in a decommissioned state.

In Azure Synapse when you scale down Apache spark pools Nodes to be scaled down will be put in a decommissioned state.