Knowledge check

Total points 5

1. Question 1

Apache Spark is a unified processing engine that can analyze big data with which of the following features?

Select all that apply.

1 / 1 point

- [x] SQL

Correct

Feedback: Spark is a unified processing engine that can analyze big data using SQL.

- [ ] Support for multiple Drivers running in parallel on a cluster
- [x] Machine Learning

Correct

Spark is a unified processing engine that can analyze big data using machine learning.

- [x] Graph Processing

Correct

Spark is a unified processing engine that can analyze big data using graph processing.

- [x] Real-time stream analysis

Correct

Spark is a unified processing engine that can analyze big data using real-time stream analysis.

2. Question 2

Which of the following Databricks features are **not** Open-Source Spark?

Select all that apply.

1 / 1 point

- [x] Databricks Workflows

Correct

Databricks Workflows is not open-source Spark.

- [x] Databricks Workspace

Correct

Databricks Workspace is not open-source Spark.

- [ ] MLFlow

- [x] Databricks Runtime

Correct

Databricks Runtime is not open-source Spark.

3. Question 3

Apache Spark notebooks allow which of the following?

Select all that apply.

1 / 1 point

- [ ] Create new Workspace
- [x] Execution of code

Correct

A notebook is a collection of cells. These cells are run to execute code.

- [x] Display graphical visualizations

Correct

A notebook is a collection of cells. These cells can display graphical visualizations.

- [x] Rendering of formatted text

Correct

A notebook is a collection of cells. These cells can be run to render formatted text.

4. Question 4

In Azure Databricks when creating a new Notebook, the default languages available to select from are?

Select all that apply.

1 / 1 point

- [ ] Java
- [x] R

Correct

In Azure Databricks when creating a new Notebook, one of the default languages available to select from is R.

- [x] SQL

Correct

In Azure Databricks when creating a new Notebook, one of the default languages available to select from is SQL.

- [x] Python

Correct

In Azure Databricks when creating a new Notebook, one of the default languages available to select from is Python.

- [x] Scala

Correct

In Azure Databricks when creating a new Notebook, one of the default languages available to select from is Scala.

5. Question 5

If your notebook is attached to a cluster, you can carry out which of the following from within the notebook?

Select all that apply.

0.75 / 1 point

- [ ] Delete the cluster
- [ ] Detach your notebook from the cluster
- [x] Attach to another cluster

Correct

If your notebook is attached to a cluster, you can attach to another cluster.

- [x] Restart the cluster

Correct

If your notebook is attached to a cluster, you can restart the cluster.

You didn't select all the correct answers

Knowledge check

Total points 6

1. Question 1

Select all that apply.

You work with Big Data as a data engineer or a data scientist, and you must process data that is oftentimes referred to as the “3 Vs of Big Data”. What do the 3Vs of Big Data stand for?

1 / 1 point

- [x] Volume

Correct

High volume - You must process an extremely large volume of data and need to scale out your compute accordingly.

- [x] Velocity

Correct

High velocity - You require streaming and real-time processing capabilities.

- [ ] Variable
- [x] Variety

Correct

Variety - Your data types are varied, from structured relational data sets and financial transactions to unstructured data such as chat and SMS messages, IoT devices, images, logs, MRIs, etc.

2. Question 2

Spark's performance is based on parallelism. Which of the following Scalability methods is limited to a finite amount of RAM, Threads and CPU speeds?

1 / 1 point

- ( ) Horizontal Scaling
- ( ) Diagonal Scaling
- (•) Vertical Scaling

Correct

Scaling vertically is limited to a finite amount of RAM, Threads and CPU speeds.

3. Question 3

In an Apache Spark Cluster jobs are divided into which of the following?

0 / 1 point

- Executors
- Slots
- Tasks
- Drivers

Incorrect

The Driver is the JVM in which our application runs.

4. Question 4

When creating a new cluster in the Azure Databricks workspace, which of the following is a sequence of steps that happens in the background?

1 / 1 point

- When an Azure Databricks workspace is deployed, you are allocated a pool of VMs. Creating a cluster draws from this pool.
- Azure Databricks creates a cluster of driver and worker nodes, based on your VM type and size selections.
- Azure Databricks provisions a dedicated VM (Virtual Machine) that processes all jobs, based on your VM type and size selection.

Correct

At the time of cluster creation, you specify the types and sizes of the virtual machines (VMs) to use for both the Driver and Worker nodes, but Azure Databricks manages all other aspects of the cluster.

5. Question 5

To parallelize work, the unit of distribution is a Spark Cluster. Every Cluster has a Driver and one or more executors. Work submitted to the Cluster is split into what type of object?

1 / 1 point

- Stages
- Arrays
- Jobs

Correct

Each parallelized action is referred to as a Job. The result of each Job is returned to the Driver. Depending on the work required, multiple Jobs will be required. Each Job is broken down into Stages.

6. Question 6

Spark Cluster use two levels of parallelization. Which of the following are levels of parallelization?

1 / 1 point

- [x] Executor

Correct

The first level of parallelization is the Executor - a Java virtual machine running on a node, typically, one instance per node.

- [ ] Partition
- [ ] Job
- [x] Slot

Correct

The second level of parallelization is the Slot - the number of which is determined by the number of cores and CPUs of each node.

Knowledge check

Total points 4

1. Question 1

How do you list files in DBFS within a notebook?

1 / 1 point

- (•) %fs ls /my-file-path
- ( ) ls /my-file-path
- ( ) %fs dir /my-file-path

Correct

Correct. You added the file system magic to the cell before executing the ls command.

2. Question 2

How do you infer the data types and column names when you read a JSON file?

1 / 1 point

- spark.read.option("inferData", "true").json(jsonFile)
- spark.read.inferSchema("true").json(jsonFile)
- spark.read.option("inferSchema", "true").json(jsonFile)

Correct

This approach is the correct way to infer the file's schema.

3. Question 3

Which of the following SparkSession functions returns a DataFrameReader

1 / 1 point

- readStream(..)
- read(..)
- emptyDataFrame(..)
- createDataFrame(..)

Correct

The function SparkSession.read() returns a DataFrameReader.

4. Question 4

When using a notebook and a spark session. We can read a CSV file. Which of the following can be used to view the first couple thousand characters of a file?

1 / 1 point

- %fs head /mnt/training/wikipedia/pageviews/pageviews_by_second.tsv
- %fs ls /mnt/training/wikipedia/pageviews/
- %fs dir /mnt/training/wikipedia/pageviews/

Correct

We can use %fs head ... to view the first couple thousand characters of a file.

Knowledge check

Total points 6

1. Question 1

Which of the following SparkSession functions returns a DataFrameReader

1 / 1 point

- createDataFrame(..)
- emptyDataFrame(..)
- .read(..)
- .readStream(..)

Correct

The function SparkSession.read() returns a DataFrameReader

2. Question 2

When using a notebook and a spark session. We can read a CSV file. Which of the following can be used to view the first couple of thousand characters of a file

1 / 1 point

- %fs head /mnt/training/wikipedia/pageviews/pageviews_by_second.tsv
- %fs ls /mnt/training/wikipedia/pageviews/
- %fs dir /mnt/training/wikipedia/pageviews/

Correct

We can use %fs head ... to view the first couple thousand characters of a file

3. Question 3

Which DataFrame method do you use to create a temporary view?

1 / 1 point

- createTempViewDF()
- createTempView()
- createOrReplaceTempView()

Correct

You use this method to create temporary views in DataFrames.

4. Question 4

How do you define a DataFrame object?

1 / 1 point

- Introduce a variable name and equate it to something like myDataFrameDF =
- Use the DF.create() syntax
- Use the createDataFrame() function

Correct

This approach is the correct way to create DataFrame objects.

5. Question 5

How do you cache data into the memory of the local executor for instant access?

1 / 1 point

- .cache()
- .inMemory().save()
- .save().inMemory()

Correct

The cache() method is an alias for persist(). Calling this moves data into the memory of the local executor.

6. Question 6

What is the Python syntax for defining a DataFrame in Spark from an existing Parquet file in DBFS?

1 / 1 point

- IPGeocodeDF = parquet.read("dbfs:/mnt/training/ip-geocode.parquet")
- IPGeocodeDF = spark.read.parquet("dbfs:/mnt/training/ip-geocode.parquet")
- IPGeocodeDF = spark.parquet.read("dbfs:/mnt/training/ip-geocode.parquet")

Correct

This syntax is correct.

Knowledge check

Total points 6

1. Question 1

Among the most powerful components of Spark are Spark SQL. At its core lies the Catalyst optimizer. When you execute code, Spark SQL uses Catalyst's general tree transformation framework in four phases.

In which order are these phases carried out?

1 / 1 point

1. logical plan optimization

2. physical planning

3. analyzing a logical plan to resolve references

4. code generation to compile parts of the query to Java bytecode

1: logical plan optimization

2: analyzing a logical plan to resolve references

3: code generation to compile parts of the query to Java bytecode

4: physical planning

1: analyzing a logical plan to resolve references

2. logical plan optimization

3: physical planning

4. code generation to compile parts of the query to Java bytecode

1: code generation to compile parts of the query to Java bytecode

2: analyzing a logical plan to resolve references

3: logical plan optimization

4: physical planning

Correct

That is the correct order

2. Question 2

Which of the following statements describes a wide transformation?

1 / 1 point

A wide transformation can be applied per partition/worker with no need to share or shuffle data to other workers

A wide transformation applies data transformation over a large number of columns

A wide transformation requires sharing data across workers. It does so by shuffling data.

Correct

Wide transformation shares data across workers by shuffling data between executors.

3. Question 3

Which of the following statements describes a narrow transformation?

1 / 1 point

Can be applied per partition/worker with no need to share or shuffle data to other workers

Requires sharing data across workers and by shuffling data.

Applies data transformation over a large number of columns

Correct

narrow transformation can be applied per partition/worker with no need to share or shuffle data to other workers.

4. Question 4

Which feature of Spark determines how your code is executed?

1 / 1 point

Catalyst Optimizer

- Tungsten Record Format
- Java Garbage Collection

Correct

Correct. Spark SQL uses Catalyst's general tree transformation framework in four phases - Analysis, Logical Optimization, Physical Planning, and Code Generation.

5. Question 5

Which feature of Spark of optimization is used in shuffling operations during wide transformations?

1 / 1 point

- Lazy Execution
- Tungsten Record Format
- Catalyst Optimizer

Correct

The Tungsten Record Format is an optimization used in shuffling operations during wide transformations. This format prevents the need for expensive serialization and de-serialization of objects in order to get data from one JVM to another.

6. Question 6

If you create a DataFrame that will read some data from Azure Blob Storage, and then you create another DataFrame by filtering the initial DataFrame. What feature of Spark causes these transformations to be analyzed?

1 / 1 point

- Tungsten Record Format
- Java Garbage Collection
- Lazy Execution

Correct

Transformations applied to DataFrames are lazy, meaning they will not trigger any jobs. If you pass the DataFrame to a display function, a job will be triggered because display is an action.

Knowledge check

Total points 6

1. Question 1

Which of the following formats are supported when importing files into an Azure Databricks notebook,?

Select all that apply.

1 / 1 point

- [x] .scala

Correct

.scala is a valid format

- [x] .Zip

Correct

.zip is a valid format

- [ ] .ORC

- [x] .dbc

Correct

.dbc is a valid format

- [x] .html

Correct

.html is a valid format

- [ ] .Yaml

2. Question 2

Examine the following code. From the options below select the correct syntax to **complete line 4** which will return an instance of a DataFrame in a Spark notebook in Azure Databricks.

1: pagecountsEnAllDF = (spark

2: .read

3: _____ # Returns an instance of DataFrame

4: .cache()

5:)

6: print(pagecountsEnAllDF)

1 / 1 point

.parquet(parquetFile)

.cache(parquetFile)

.DataFrame(parquetFile)

.read(parquetFile)

Correct

.parquet(parquetFile) can be used to return an instance of a DataFrame

3. Question 3

Examine the following piece of code taken from a notebook in an Azure Databricks.

Complete line 4 so that 15 rows of data will be displayed, and the columns will not be truncated.

1: sortedDF = (pagecountsEnAllDF

2: .orderBy("requests")

3:

4: SortedDF. _____

0 / 1 point

sortedDF.print(15)

sortedDF.print(15, False)

sortedDF.show(15)

sortedDF.show(15, False)

Incorrect

This will sort but will truncate columns

4. Question 4

Which command will order by a column in descending order?

1 / 1 point

df.orderBy("requests").desc()

df.orderBy(col("requests").desc())

df.orderBy("requests desc")

Correct

Use the desc() method on the Column Class to reverse the order.

5. Question 5

Which command specifies a column value in a DataFrame's filter? Specifically, filter by a productType column where the value is equal to book?

1 / 1 point

- df.col("productType").filter("book")

This syntax is incorrect. There is no col method on a DataFrame.

- df.filter(col("productType") == "book")
- df.filter("productType = 'book'")

Correct

This approach is the correct way to apply the filter, by using the Column Class

6. Question 6

When using the Column Class, which command filters based on the end of a column value? For example, a column named verb and filtered by words ending with "ing".

1 / 1 point

- df.filter("verb like '%ing'")
- df.filter(col("verb").endswith("ing"))
- df.filter().col("verb").like("%ing")

Correct

The Column Class supports both the endswith() method and the like() method (example - col("verb").like("%ing")).

Knowledge check

Total points 5

1. Question 1

Which of the listed methods for renaming a DataFrame's column are correct?

Select two options.

1 / 1 point

- [x] C: df.toDF("dateCaptured")

Correct

This is a valid renaming method.

- [x] df.select(col("timestamp").alias("dateCaptured"))

Correct

Feedback: This is a valid renaming method.

- [ ] df.alias("timestamp", "dateCaptured")

2. Question 2

You need to find the average of sales transactions by storefront. Which of the following aggregates would you use?

1 / 1 point

- (•) df.groupBy(col("storefront")).avg("completedTransactions")
- ( ) df.groupBy(col("storefront")).avg(col("completedTransactions"))
- ( ) df.select(col("storefront")).avg("completedTransactions")

Correct

Feedback: The syntax shown groups the data by the storefront Column, then calculates the average value of completed sales transactions.

3. Question 3

In Azure Databricks you are about to do some ETL on a file you have received from a customer. The file contains data about people, including:

first, middle, and last names

gender

birth date

Social Security number

Salary

You discover that the file contains some duplicate records and you have been instructed to remove any duplicates. The dropDuplicates() command will more than likely create a shuffle. To help reduce the number of post-shuffle partitions which of the following commands should you run?

1 / 1 point

- spark.sql.conf.set("spark.shuffle.partitions", 8)
- spark.conf.set("spark.sql.partitions", 8)
- spark.conf.set("spark.sql.shuffle.partitions", 8)

Correct

Feedback: spark.conf.set("spark.sql.shuffle.partitions", 8) is the correct syntax.

4. Question 4

Which of the following syntax will successfully display the year portion for a column named capturedAt and formatted as a Timestamp column?

1 / 1 point

- .select(year(col("capturedAt")))
- .select(year ("capturedAt")
- .select(col("capturedAt")year)

Correct

This is the correct syntax to return the year portion of a Timestamp formatted column.

5. Question 5

You need to change a column name from “dob” to “DateOfBirth” on a spark DataFrame. Which of the following syntax is valid?

1 / 1 point

- .ColumnRename("dob","DateOfBirth")
- .RenameColumn("dob","DateOfBirth")
- .withColumnRenamed("dob","DateOfBirth")

Correct

This is correct and will rename the column “dob” to “DateOfBirth”.

Knowledge check

Total points 6

1. Question 1

True or False?

ETL/ELT workflows including analytics workloads in Azure Databricks can be operationalized using Azure Data Factory pipelines.

1 / 1 point

- [x] True
- [ ] False

Correct

ETL/ELT workflows (including analytics workloads in Azure Databricks) can be operationalized using Azure Data Factory pipelines.

2. Question 2

When you create an Azure Databricks service, a "Databricks appliance" is deployed as an Azure resource in your subscription. When a Databricks appliance is deployed into Azure which of the following resources are created?

Select all that apply.

1 / 1 point

- [x] Virtual Network

Correct

A Virtual Network is deployed.

- [x] Network security Group

Correct

an NSG is deployed.

- [ ] Azure SQL Database

3. Question 3

In Azure Data Bricks the Blob Storage account provides default file storage within the workspace referred to as DBFS.

What does DBFS stand for?

0 / 1 point

- Data Block File System
- Database File system
- Databricks File System

Incorrect

DBFS does not stand for Database File system.

4. Question 4

In Azure Databricks when ADLS Passthrough is configured on a standard cluster you must set which of the following?

1 / 1 point

- Group Access
- Single User Access
- Multiple Users

Correct

On a standard cluster, when you enable this setting, you must set single user access to one of the Azure Active Directory (AAD) users in the Azure Databricks workspace.

5. Question 5

By default, all users can create and modify clusters unless an administrator enables cluster access control. With cluster access control, permissions determine a user's abilities. There are four permission levels for a cluster. **Select the correct four permissions.**

1 / 1 point

- [x] Can Attach To

Correct

Can Attach To is a valid permission level for a cluster.

- [ ] Can Edit
- [x] Can Restart

Correct

Can Restart is a valid permission level for a cluster.

☑ Can Manage

Correct

Can Manage is a valid permission level for a cluster.

☑ No Permissions

Correct

No Permissions is a valid permission level for a cluster.

☐ Can Read

6. Question 6

Users access Azure Databricks workspace with an Azure AD account

Is the following statement True or False?

The user's Azure AD account has to be added to the Azure Databricks workspace before they can access it.

1 / 1 point

◉ True

○ False

Correct

The user's Azure AD account has to be added to the Azure Databricks workspace before they can access it.

Knowledge check

Total points 6

1. Question 1

Delta Lake provides snapshots of data enabling developers to access and revert to earlier versions of data for audits, rollbacks or to reproduce experiments. This functionality is referred to as?

1 / 1 point

- Schema Evolution
- ACID Transactions
- Time Travel
- Schema Enforcement

Correct

Delta Lake provides snapshots of data enabling developers to access and revert to earlier versions of data for audits, rollbacks or to reproduce experiments.

2. Question 2

One of the core features of Delta Lake is performing upserts. Which of the following statements is true in regard to Upsert?

1 / 1 point

- UpSert is literally TWO operations. Update / Insert
- Upsert is supported in traditional data lakes
- Upsert is a new DML statement for SQL syntax

Correct

To UPSERT means to "UPdate" and "inSERT". In other words, UPSERT is literally TWO operations. It is not supported in traditional data lakes.

3. Question 3

When discussing Delta Lake, there is often a reference to the concept of Bronze, Silver and Gold tables. These levels refer to the state of data refinement as data flows through a processing pipeline and are conceptual guidelines. Based on these table concepts the refinements in Silver tables generally relate to **which of the following?**

1 / 1 point

- Highly refined views of the data
- Raw data (or very little processing)
- Data that is directly queryable and ready for insights

Correct

Silver tables generally relate to data that is directly queryable and ready for insights.

4. Question 4

What is the Databricks Delta command to display metadata?

1 / 1 point

- SHOW SCHEMA table name
- DESCRIBE DETAIL table Name
- MSCK DETAIL table name

Correct

You display metadata by using DESCRIBE DETAIL table Name.

5. Question 5

How do you perform UPSERT in a Delta dataset?

1 / 1 point

- Use MERGE INTO my-table USING data-to-upsert
- Use UPSERT INTO my-table
- Use UPSERT INTO my-table /MERGE

Correct

That's the correct syntax to perform UPSERT in a Databricks Delta dataset.

6. Question 6

What optimization does the following command perform: OPTIMIZE Students ZORDER BY Grade?

1 / 1 point

- Ensures that all data backing, for example, Grade=8 is colocated, then rewrites the sorted data into new Parquet files.
- Creates an order-based index on the Grade field to improve filters against that field.
- Ensures that all data backing, for example, Grade=8 is colocated, then updates a graph that routes requests to the appropriate files.

Correct

ZOrdering colocates related information in the same set of files.

Knowledge check

Total points 6

1. Question 1

The lambda architecture is a big data processing architecture combining both batch and real-time processing methods and features an append-only immutable data source. Which of the following are features of an append-only immutable data source?

Select all that apply.

1 / 1 point

- [x] Data is implicitly ordered by time of arrival

Correct

Data is implicitly ordered by time of arrival.

- [x] serves as system of record

Correct

Lambda features an append-only immutable data source that serves as system of record. A system of record (SOR) is an information storage and retrieval system that can serve as an authoritative source of truth,

- [x] Timestamped events are appended to existing events

Correct

Timestamped events are appended to existing events (nothing is overwritten).

- [ ] Timestamped events overwrite existing events

2. Question 2

Delta Lake Architecture improves upon the traditional Lambda architecture through a unified pipeline that allows you to combine batch and streaming workflows through a shared filestore with ACID-compliant transactions. **What do the letters ACID stand for?**

Select 4 options.

1 / 1 point

- [x] Isolation

Correct

Isolation ensures that concurrent execution of transactions leaves the database in the same state that would have been obtained if the transactions were executed sequentially.

- [x] Consistency

Correct

Consistency ensures that a transaction can only bring the database from one valid state to another.

- [ ] Desirable
- [ ] Implicit
- [ ] Concurrency
- [x] Atomicity

Correct

Atomicity guarantees that each transaction is treated as a single unit which either succeeds completely or fails completely.

- [ ] Agile
- [x] Durability

Correct

Durability guarantees that once a transaction has been committed, it will remain committed even in the case of a system failure.

3. Question 3

In the Delta Lake architecture, the refinement of the data is often referred to as Bronze, Silver and Gold Tables. Which of the following tables provide business level aggregates often used for reporting and Dashboarding?

1 / 1 point

- ( ) Bronze
- ( ) Silver
- (•) Gold

Correct

Gold tables provide business level aggregates often used for reporting and dashboarding. This would include aggregations such as daily active website users, weekly sales per store, or gross revenue per quarter by department.

4. Question 4

What is a lambda architecture and what does it try to solve?

1 / 1 point

An architecture that defines a data processing pipeline whereby microservices act as compute resources for efficient large-scale data processing.

An architecture that employs the latest Scala runtimes in one or more Databricks clusters to provide the most efficient data processing platform available today.

An architecture that splits incoming data into two paths - a batch path and a streaming path. This architecture helps address the need to provide real-time processing in addition to slower batch computations.

Correct

The lambda architecture is a big data processing architecture that combines both batch- and real-time processing methods.

5. Question 5

What command should be issued to view the list of active streams?

1 / 1 point

Invoke spark.streams.show

Invoke spark.streams.active

Invoke spark.view.active

Correct

That's the correct syntax to view the list of active streams.

6. Question 6

What is required to specify the location of a checkpoint directory when defining a Delta Lake streaming query?

1 / 1 point

.writeStream.format("delta").option("checkpointLocation", checkpointPath) ...

.writeStream.format("delta").checkpoint("location", checkpointPath) ...

.writeStream.format("parquet").option("checkpointLocation", checkpointPath) ...

Correct

That's the correct syntax to specify the checkpoint directory on a Delta Lake streaming query.

Knowledge check

Total points 6

1. Question 1

Stream processing is where you continuously incorporate new data into Data Lake storage and compute results. Which of the following would be examples of Stream processing?

1 / 1 point

- [ ] Invoicing
- [x] Bank Card Processing

Correct

A stream of data is treated as a table to which data is continuously appended Bank Card Processing would be an example of stream processing.

- [ ] Monthly Payroll processing
- [x] Game play events

Correct

A stream of data is treated as a table to which data is continuously appended Game play events would be an example of stream processing.

- [x] IoT Device Data

Correct

A stream of data is treated as a table to which data is continuously appended IoT Device Data would be an example of stream processing.

2. Question 2

When creating a new event hub in the Azure Portal you are required to specify a Namespace name. The namespace name must be unique in which of the following?

1 / 1 point

- Resource group only
- Tenant Only

Azure

Subscription only

Correct

An Event Hubs namespace provides a unique scoping container, in which you create one or more event hubs.

3. Question 3

When doing a write stream command, what does the outputMode("append") option do?

1 / 1 point

The append mode allows records to be updated and changed in place

The append outputMode allows records to be added to the output sink

The append mode replaces existing records and updates aggregates

Correct

The outputMode "append" option informs the write stream to add only new records to the output sink. The "complete" option is to rewrite the full output - applicable to aggregations operations. Finally, the "update" option is for updating changed records in place.

4. Question 4

In Spark Structured Streaming, what method should be used to read streaming data into a DataFrame?

1 / 1 point

spark.readStream

spark.read

spark.stream.read

Correct

Use the spark.readStream method to start reading data from a streaming query into a DataFrame.

5. Question 5

What happens if the command option("checkpointLocation", pointer-to-checkpoint directory) is not specified?

1 / 1 point

The streaming job will function as expected since the checkpointLocation option does not exist

When the streaming job stops, all state around the streaming job is lost, and upon restart, the job must start from scratch.

It will not be possible to create more than one streaming query that uses the same streaming source since they will conflict

Correct

Setting the checkpointLocation is required for many sinks used in Structured Streaming. For those sinks where this setting is optional, keep in mind that when you do not set this value, you risk losing your place in the stream.

6. Question 6

Select the correct option to complete the statement below:

In Azure Databricks every streaming DataFrame must have a schema. That is the definition of column names and data types. For file based streaming sources the schema is _____.

1 / 1 point

Defined for you

User Defined

Both Defined for you and can be user defined if required

Correct

For file-based streaming sources, the schema must be user-defined.

Knowledge check

Total points 6

1. Question 1

In Azure Databricks when creating a new user access token, the Lifetime setting of the access token can be manually set. What is the default Lifetime (Days) value when creating a new access token?

0 / 1 point

- 120 days
- 30 Days
- 60 Days
- 90

Incorrect

That is not the default lifetime (Days).

2. Question 2

In Azure Databricks when creating a new user access token, the Lifetime setting of the access token can be manually set. If the Token Lifetime is unspecified what will be the Lifetime(Days) of the token?

1 / 1 point

- Indefinite
- 60 Days
- 30 Days
- 120 Days
- 90 Days

Correct

If the lifetime is unspecified then the token will have an indefinite lifetime (Days).

3. Question 3

True or False?

In Azure Databricks, personal access tokens can be used for secure authentication to the Databricks API instead of passwords. After a new token is generated, it can be viewed by going back to the user settings from where it was generated.

1 / 1 point

- False

True

Correct

Feedback: When generating a new token, you will be presented with the unique token and advised to copy the token as it will not be visible again once created.

4. Question 4

What's the purpose of linked services in Azure Data Factory?

1 / 1 point

To link data stores or computer resources together for the movement of data between resources

To represent a processing step in a pipeline

To represent a data store or a compute resource that can host execution of an activity

Correct

Linked services define the connection information needed for Data Factory to connect to external resources.

5. Question 5

How can parameters be passed into an Azure Databricks notebook from Azure Data Factory?

1 / 1 point

Deploy the notebook as a web service in Databricks, defining parameter names and types

Use the new API endpoint option on a notebook in Databricks and provide the parameter name

Use notebook widgets to define parameters that can be passed into the notebook

Correct

You can configure parameters by using widgets on the Databricks notebook. You then pass in parameters with those names via a Databricks notebook activity in Data Factory.

6. Question 6

What happens to Databricks activities (notebook, JAR, Python) in Azure Data Factory if the target cluster in Azure Databricks isn't running when the cluster is called by Data Factory?

1 / 1 point

Simply add a Databricks cluster start activity before the notebook, JAR, or Python Databricks activity

The Databricks activity will fail in Azure Data Factory – you must always have the cluster running

If the target cluster is stopped, Databricks will start the cluster before attempting to execute

Correct

This situation will result in a longer execution time because the cluster must start, but the activity will still execute as expected.

Knowledge check

Total points 6

1. Question 1

Azure DevOps is a collection of services that provide an end-to-end solution for the five core practices of DevOps. The **five** core practices of DevOps as defined by Microsoft are?

1 / 1 point

- [x] Development

Correct

is defined as a core practice in DevOps.

- [ ] Scoping
- [ ] Program Management
- [x] Monitoring and Operations

Correct

Monitoring and Operations is defined as a core practice in DevOps

- [x] Planning and Tracking

Correct

Planning and Tracking is defined as a core practice in DevOps.

- [ ] Project Management
- [x] Delivery

Correct

is defined as a core practice in DevOps

- [x] Build and Test

Correct

Build and Test is defined as a core practice in DevOps.

2. Question 2

In an Azure DevOps project creating a release pipeline provides which of the following portions of CI/CD?

1 / 1 point

- (•) CD
- ( ) CI

Correct

A release pipeline provides the CD portion of CI/CD.

3. Question 3

What does the CD in CI/CD mean?

1 / 1 point

- ( ) Continuous Delivery
- (•) Both are correct
- ( ) Continuous Deployment

Correct

Continuous Delivery automates your release process up to the point where human intervention is needed, by clicking a button. Continuous Deployment takes a step further by removing the human intervention and relying on automated tests to automatically determine whether the build should be deployed into production.

4. Question 4

What sort of pipeline is required in Azure DevOps for creating artifacts used in releases?

1 / 1 point

- ( ) An Artifact pipeline

A Build pipeline

A Release pipeline

Correct

The output of a Build pipeline is one or more artifacts that can be used within release pipelines for automated deployments.

5. Question 5

What steps are required to authorize Azure DevOps to connect to and deploy notebooks to a staging or production Azure Databricks workspace?

1 / 1 point

Create an Azure Active Directory application, copy the application ID, then use that as the Databricks bearer token in the Databricks Notebooks Deployment step of the Release pipeline

In the production or staging Azure Databricks workspace, enable Git integration to Azure DevOps, then link to the Azure DevOps source code repo

Create a new Access Token within the user settings in the production Azure Databricks workspace, then use the token as the Databricks bearer token in the Databricks Notebooks Deployment step of the Release pipeline

Correct

The Access Token allows you to grant access to resources within an Azure Databricks workspace without passing in user credentials.

6. Question 6

In an Azure DevOps project creating a build pipeline provides which of the following portions of CI/CD

1 / 1 point

CI

CD

Correct

A Build pipeline provides the CI portion of CI/CD

Knowledge check

Total points 5

1. Question 1

What are the two prerequisites for connecting Azure Databricks with Azure Synapse Analytics that apply to the Azure Synapse Analytics instance?

1 / 1 point

- Add the client IP address to the firewall's allowed IP addresses list and use the correctly formatted ConnectionString
- Create a database master key and configure the firewall to enable Azure services to connect
- Use a correctly formatted ConnectionString and create a database master key

Correct

Create a database master key and configure the firewall to enable Azure services to connect.

2. Question 2

Which of the following is the correct syntax for overwriting data in Azure Synapse Analytics from a Databricks notebook?

1 / 1 point

- df.write.format("com.databricks.spark.sqldw").mode("overwrite").option("...").option("...").save()
- df.write.mode("overwrite").option("...").option("...").save()
- df.write.format("com.databricks.spark.sqldw").overwrite().option("...").option("...").save()

Correct

The key is to specify the correct format, intended write mode, and options that specify the Azure Synapse Analytics properties.

3. Question 3

The Azure Synapse Connector uses Azure Blob Storage as intermediary storage and using PolyBase in Synapse enables MPP reads and writes to Synapse from Azure Databricks. However, the Synapse connector is more suited to ETL than to interactive queries. For interactive and ad-hoc queries, data should be extracted into which of the following?

1 / 1 point

- Azure SQL database Table
- Azure Databricks Delta table
- Azure Data Factory table

Correct

For interactive and ad-hoc queries, data should be extracted into a Databricks Delta table.

4. Question 4

You can access Azure Synapse from Databricks using the Azure Synapse connector which uses three types of network connections.

Which of the following connections are used by Synapse Connector?

Select all that apply.

1 / 1 point

- [x] Spark driver to Azure Synapse

Correct

Spark driver to Azure Synapse is one of the 3 connections used by Synapse connector.

- [x] Spark driver and executors to Azure storage account

Correct

Spark driver and executors to Azure storage account is one of the 3 connections used by Synapse connector.

- [x] Azure Synapse to Azure storage account

Correct

Azure Synapse to Azure storage account is one of the 3 connections used by Synapse connector.

- [ ] Azure Storage account to Databricks connection
- [ ] Databricks connector to Spark driver

5. Question 5

You can access Azure Synapse from Databricks using the Azure Synapse connector and it is recommended that the connection strings use Secure Sockets Layer (SSL) encryption for all data sent between the Spark driver and the Azure Synapse instance through the JDBC connection. To verify that SSL encryption is enabled, you should verify that which of the following is set in the connection string?

1 / 1 point

- ( ) encrypt=enabled
- (•) encrypt=true
- ( ) encrypt=on
- ( ) encrypt=active

Correct

To verify that SSL encryption is enabled, you can search for encrypt=true in the connection string.

Knowledge check

Total points 6

1. Question 1

Select two items from the following options to complete this statement correctly:

Azure Databricks uses Azure Active Directory (AAD) as the exclusive Identity Provider. Any AAD member assigned to the _____ or _____ role can deploy Databricks and is automatically added to the ADB members list upon first login.

1 / 1 point

- [x] Contributor

Correct

Any AAD member assigned to the Contributor role can deploy Databricks.

- [ ] Reader
- [ ] User Access Administrator
- [x] Owner

Correct

Any AAD member assigned to the Owner role can deploy Databricks

2. Question 2

Azure Databricks is a multitenant service and to provide fair resource sharing to all regional customers, it imposes limits on API calls. What is currently the restrictions on the maximum number of notebooks or execution contexts that can be attached to a cluster?

1 / 1 point

- 100
- No Limit
- 200
- 150

Correct

There can be a maximum of 150 notebooks or execution contexts attached to a cluster.

3. Question 3

Azure Databricks deployments are built on top of the Azure infrastructure and currently have default restrictions or Azure limits. Currently, what is the maximum number of storage accounts per region per subscription in Azure Databricks?

1 / 1 point

- 500
- 1000
- 250
- 150

Correct

250 is the maximum number of storage accounts per region per subscription.

4. Question 4

Azure Databrick jobs use clusters and different types of jobs demand different types of cluster resources. When training machine learning models you should consider using which of the following?

1 / 1 point

- General purpose VMs
- Computing optimized VMs
- Autoscaling features
- Memory optimized VMs

Correct

To train machine learning models its required cache all of the data in memory. Consider using memory optimized VMs so that the cluster can take advantage of the RAM cache.

5. Question 5

What is SCIM?

1 / 1 point

- (●) An open standard that enables organizations to import both groups and users from Azure Active Directory into Azure Databricks
- ( ) An open standard that enables users to bring their own auth key to the Databricks environment
- ( ) An optimization that removes orphaned data from a given dataset

Correct

By default, Azure Active Directory roles have no relationship with groups created inside of Azure Databricks. SCIM enables synchronizing users and groups, and synchronization is automatic after initial import.

6. Question 6

If mounting an Azure Data Lake Storage (ADLS) account to a workspace, what cluster feature must be used to have ACLs within ADLS applied to the user executing commands in a notebook?

1 / 1 point

- (●) Enable ADLS Passthrough on a cluster.
- ( ) Set spark.config.adls.impersonateuser(true)
- ( ) Enable SCIM

Correct

When enabled, authentication automatically takes place in Azure Data Lake Storage (ADLS) from Azure Databricks clusters using the same Azure Active Directory (Azure AD) identity that one uses to log into Azure Databricks. Any ACLs applied at the folder or file level in ADLS are enforced based on the user's identity.

Test prep

Latest Submission Grade 87.5%

1. Question 1

Azure Databricks Runtime adds several key capabilities to Apache Spark workloads that can increase performance and reduce costs. Which of the following are features of Azure Databricks?

Select all that apply.

1 / 1 point

- [x] High-speed connectors to Azure storage services

Correct

Azure Databricks Runtime adds several key capabilities to Apache Spark workloads that can increase performance and reduce costs including High-speed connectors to Azure storage services.

- [ ] Parallel Cluster Drivers
- [x] Indexing

Correct

Azure Databricks Runtime adds several key capabilities to Apache Spark workloads that can increase performance and reduce costs including Indexing.

- [x] Caching

Correct

Azure Databricks Runtime adds several key capabilities to Apache Spark workloads that can increase performance and reduce costs including Caching.

- [x] Auto-scaling and auto-termination

Correct

Try going back and reviewing the Describe Azure Databricks lesson.

2. Question 2

Apache Spark supports which of the following languages?

Select all that apply.

0.75 / 1 point

- [x] Scala

Correct

Apache Spark supports Scala.

- [ ] ORC

- [x] Java
- [x] Python

Correct

Apache Spark supports Python.
You didn't select all the correct answers

3. Question 3

Which of the following statements are True

Select all that apply.

1 / 1 point

- [x] To use your Azure Databricks notebook to run code, you must attach it to a cluster

Correct

To use your notebook to run a code, you must attach it to a cluster.

- [x] You can detach a notebook from a cluster and attach it to another cluster.

Correct

detach your notebook from a cluster and attach it to another depending upon your organization's requirements.

- [ ] Once created a notebook can only be connected to the original cluster.
- [ ] To use your Azure Databricks notebook to run code you do not require a cluster

4. Question 4

Which of the following Databricks features are not Open-Source Spark?

1 / 1 point

- [ ] MLFlow
- [x] Databricks Runtime

Correct

Databricks Runtime is not open-source Spark

- [x] Databricks Workflows

Correct

Databricks Workflows is not open-source Spark

- [x] Databricks Workspace

Correct

Databricks Workspace is not open-source Spark

5. Question 5

How many drivers does a Cluster have?

1 / 1 point

- Configurable between one and eight
- Only one
- Two, running in parallel

Correct

Feedback: A Cluster has one and only one driver.

6. Question 6

What type of process are the driver and the executors?

1 / 1 point

- C++ processes
- Python processes
- Java processes

Correct

The driver and the executors are Java processes.

7. Question 7

You work with Big Data as a data engineer, and you must process real-time data. This is referred to as having which of the following characteristics?

1 / 1 point

- High velocity
- High volume
- Variety

Correct

This characteristic relates to the requirement for streaming and real-time processing capabilities.

8. Question 8

Spark's performance is based on parallelism. Which of the following Scalability methods is limited to a finite amount of RAM, Threads and CPU speeds?

1 / 1 point

- Diagonal Scaling

- [x] Vertical Scaling
- [ ] Horizontal Scaling

Correct

Scaling vertically is limited to a finite amount of RAM, Threads and CPU speeds.

9. Question 9

Spark Cluster use two levels of parallelization. Which of the following are levels of parallelization?

1 / 1 point

- [x] Executor

Correct

The first level of parallelization is the Executor - a Java virtual machine running on a node, typically, one instance per node.

- [ ] Job
- [ ] Partition
- [x] Slot

Correct

The second level of parallelization is the Slot - the number of which is determined by the number of cores and CPUs of each node.

10. Question 10

In an Apache Spark Cluster jobs are divided into which of the following?

0 / 1 point

- [ ] Tasks
- [x] Slots
- [ ] Drivers
- [ ] Executors

Incorrect

Try going back and reviewing the Spark architecture fundamentals lesson.

Test prep

Latest Submission Grade 83.33%

1. Question 1

How do you list files in DBFS within a notebook?

1 / 1 point

- %fs dir /my-file-path
- %fs ls /my-file-path
- ls /my-file-path

Correct

Feedback: Correct. You added the file system magic to the cell before executing the ls command.

2. Question 2

How do you infer the data types and column names when you read a JSON file?

1 / 1 point

- spark.read.option("inferSchema", "true").json(jsonFile)
- spark.read.inferSchema("true").json(jsonFile)
- spark.read.option("inferData", "true").json(jsonFile)

Correct

This approach is the correct way to infer the file's schema.

3.

Question 3

Which of the following SparkSession functions returns a DataFrameReader?

1 / 1 point

- readStream(..)
- read(..)
- createDataFrame(..)
- emptyDataFrame(..)

Correct

The function SparkSession.read() returns a DataFrameReader.

4. Question 4

When using a notebook and a spark session. We can read a CSV file. Which of the following can be used to view the first couple thousand characters of a file?

1 / 1 point

- %fs dir /mnt/training/wikipedia/pageviews/
- %fs ls /mnt/training/wikipedia/pageviews/
- %fs head /mnt/training/wikipedia/pageviews/pageviews_by_second.tsv

Correct

We can use %fs head ... to view the first couple thousand characters of a file.

5. Question 5

You have created an Azure Databricks cluster, and you have access to a source file.

fileName = "dbfs:/mnt/training/wikipedia/clickstream/2015_02_clickstream.tsv"

You need to determine the structure of the file. Which of the following commands will assist with determining what the column and data types are?

1 / 1 point

- .option("header", "false")
- .option("header", "true")
- .option("inferSchema", "true")
- .option("inferSchema", "false")

Correct

using .option("inferSchema", "true") Spark will automatically go through the file and infer the schema of each column.

6. Question 6

In an Azure Databricks workspace, you run the following command:

%fs head /mnt/training/wikipedia/pageviews/pageviews_by_second.tsv

The partial output from this command is as follows:

[Truncated to first 65536 bytes]

"timestamp" "site" "requests"

"2015-03-16T00:09:55" "mobile" 1595

"2015-03-16T00:10:39" "mobile" 1544

"2015-03-16T00:19:39" "desktop" 2460

"2015-03-16T00:38:11" "desktop" 2237

"2015-03-16T00:42:40" "mobile" 1656

"2015-03-16T00:52:24" "desktop" 2452

Which of the following pieces of information can be inferred from the command and the output?

Select all that apply.

0.6666666666666666 / 1 point

- [x] The column is Tab separated

Correct

Feedback: The file is tab separated. This can be inferred from the file extension and the lack of other characters between each "column".

- [ ] The file has no header

- [x] Two columns are strings, and one column is a number

Correct

The strings are enclosed in double quotes while the number column is not

- [ ] The file has a header

- [x] All columns are strings

This should not be selected

Try going back and reviewing the Use Azure Databricks to prepare the data for advanced analytics and machine learning operations lesson.

- [ ] the file is a comma separated or CSV file

7. Question 7

In an Azure Databricks you wish to create a temporary view that will be accessible to multiple notebooks. Which of the following commands will provide this feature?

0 / 1 point

- createOrReplaceGlobalTempView(..)

- createOrReplaceTempView(..)

- createOrReplaceTempView(set_scope “Global”)

Incorrect

Try going back and reviewing the Use Azure Databricks to prepare the data for advanced analytics and machine learning operations lesson.

8. Question 8

Which of the following is true in respect of Parquet Files?

Select all that apply.

1 / 1 point

- [x] D: Is a splittable "file format".

Correct

Parquet files are splittable.

- [x] E: Is a Column-Oriented data store

Correct

Parquet files are Column-Oriented.

- [x] Open Source

Correct

Parquet files are free Open Source.

- [ ] Designed for performance on small data sets
- [ ] Is a Row-Oriented data store
- [x] Efficient data compression

Correct

Parquet files provide efficient data compression.

Test prep

Latest Submission Grade 100%

1. Question 1

When creating a new cluster in Azure Databricks there are three Cluster Modes that can be set. Which of the following are valid Cluster Modes?

Select three valid options.

1 / 1 point

- [x] Single Node

Correct

Single Node is a valid option.

- [ ] Multi Node
- [ ] Low Concurrency
- [x] High Concurrency

Correct

High Concurrency is a valid option.

- [x] Standard

Correct

Standard is a valid Cluster Mode.

2. Question 2

Which DataFrame method do you use to create a temporary view?

1 / 1 point

- (•) createOrReplaceTempView()
- ( ) createTempViewDF()
- ( ) createTempView()

Correct

Feedback: You use this method to create temporary views in DataFrames.

3. Question 3

How do you define a DataFrame object?

1 / 1 point

- ( ) Use the DF.create() syntax

- [x] Introduce a variable name and equate it to something like myDataFrameDF =
- [ ] Use the createDataFrame() function

Correct

This approach is the correct way to create DataFrame objects.

4. Question 4

How do you cache data into the memory of the local executor for instant access?

1 / 1 point

- [ ] .save().inMemory()
- [ ] .inMemory().save()
- [x] .cache()

Correct

Feedback: The cache() method is an alias for persist(). Calling this moves data into the memory of the local executor.

5. Question 5

What is the Python syntax for defining a DataFrame in Spark from an existing Parquet file in DBFS?

1 / 1 point

- [ ] IPGeocodeDF = spark.parquet.read("dbfs:/mnt/training/ip-geocode.parquet")
- [x] IPGeocodeDF = spark.read.parquet("dbfs:/mnt/training/ip-geocode.parquet")
- [ ] IPGeocodeDF = parquet.read("dbfs:/mnt/training/ip-geocode.parquet")

Correct

Feedback: This syntax is correct.

6. Question 6

Among the most powerful components of Spark are Spark SQL. At its core lies the Catalyst optimizer. When you execute code, Spark SQL uses Catalyst's general tree transformation framework in four phases. In which order are these phases carried out?

1: logical plan optimization

2: analyzing a logical plan to resolve references

3: code generation to compile parts of the query to Java bytecode

4: physical planning

1 / 1 point

- 3, 2, 1, 4
- 2, 3, 1, 4
- 1, 2, 3, 4
- [x] 2, 1, 4, 3

Correct

That is the correct order.

7. Question 7

Which of the following statements describes a wide transformation?

1 / 1 point

- A wide transformation applies data transformation over a large number of columns
- [x] A wide transformation requires sharing data across workers. It does so by shuffling data.
- A wide transformation can be applied per partition/worker with no need to share or shuffle data to other workers

Correct

Wide transformation shares data across workers by shuffling data between executors.

8. Question 8

Which of the following statements describes a narrow transformation?

1 / 1 point

- [x] A narrow transformation can be applied per partition/worker with no need to share or shuffle data to other workers

Correct

narrow transformation can be applied per partition/worker with no need to share or shuffle data to other workers.

- [x] A narrow transformation requires sharing data across workers. It does so by shuffling data.

Correct

this describes a Wide transformation.

- [ ] A narrow transformation applies data transformation over a large number of columns

9. Question 9

Which feature of Spark determines how your code is executed?

1 / 1 point

- Java Garbage Collection
- Tungsten Record Format
- Catalyst Optimizer

Correct

Correct. Spark SQL uses Catalyst's general tree transformation framework in four phases - Analysis, Logical Optimization, Physical Planning, and Code Generation.

10. Question 10

Which feature of Spark of optimization is used in shuffling operations during wide transformations?

1 / 1 point

- Tungsten Record Format
- Lazy Execution
- Catalyst Optimizer

Correct

The Tungsten Record Format is an optimization used in shuffling operations during wide transformations. This format prevents the need for expensive serialization and de-serialization of objects in order to get data from one JVM to another.

Test prep

Latest Submission Grade 100%

1. Question 1

When importing files in an Azure Databricks notebook, which of the following formats are supported?

Select all that apply.

1 / 1 point

- [x] .scala

Correct

.scala is a valid format.

- [x] .Zip

Correct

.zip is a valid format.

- [x] .dbc

Correct

Try going back and reviewing the Work with DataFrames columns in Azure Databricks lesson.

- [x] .html

Correct

.html is a valid format

- [ ] .ORC
- [ ] .Yaml

2. Question 2

Examine the following code. From the options below select the correct syntax **to complete line 3** so that it will return an instance of a DataFrame in a spark notebook in Azure Databricks.

1: pagecountsEnAllDF = (spark

2: .read

3: _____ # Returns an instance of DataFrame

4: .cache()

5:)

6: print(pagecountsEnAllDF)

1 / 1 point

- .DataFrame(parquetFile)
- .cache(parquetFile)
- .parquet(parquetFile)
- .read(parquetFile)

Correct

.parquet(parquetFile) can be used to return an instance of a DataFrame.

3. Question 3

Examine the following piece of code taken from a notebook in an Azure Databricks.

Complete line 4 so that 15 rows of data will be displayed, and the columns will not be truncated.

1: sortedDF = (pagecountsEnAllDF

2: .orderBy("requests")

3:)

4: SortedDF. _____

1 / 1 point

- sortedDF.print(15, False)
- sortedDF.show(15)
- sortedDF.show(15, False)
- sortedDF.print(15)

Correct

This will sort and will not truncate columns.

4. Question 4

Which command will order by a column in descending order?

1 / 1 point

- df.orderBy(col("requests").desc())
- df.orderBy("requests desc")
- df.orderBy("requests").desc()

Correct

Feedback: Use the desc() method on the Column Class to reverse the order.

5. Question 5

Which command specifies a column value in a DataFrame's filter? Specifically, filter by a productType column where the value is equal to book?

1 / 1 point

- [x] df.filter(col("productType") == "book")
- [ ] df.col("productType").filter("book")
- [ ] df.filter("productType = 'book'")

Correct

This approach is the correct way to apply the filter, by using the Column Class.

6. Question 6

When using the Column Class, which command filters are based on the end of a column value? For example, a column named verb and filtered by words ending with "ing"?

1 / 1 point

- [ ] df.filter().col("verb").like("%ing")
- [ ] df.filter("verb like '%ing'")
- [x] df.filter(col("verb").endswith("ing"))

Correct

Feedback: The Column Class supports both the endswith() method and the like() method (example - col("verb").like("%ing"))

7. Question 7

Which of the listed methods for renaming a DataFrame's column are correct?

Select two.

1 / 1 point

- [x] df.toDF("dateCaptured")

Correct

This is a valid renaming method.

- [ ] df.alias("timestamp", "dateCaptured")
- [x] df.select(col("timestamp").alias("dateCaptured"))

Correct

This is a valid renaming method.

8. Question 8

You need to find the average of sales transactions by storefront. Which of the following aggregates would you use?

1 / 1 point

- df.groupBy(col("storefront")).avg("completedTransactions")
- df.select(col("storefront")).avg("completedTransactions")
- df.groupBy(col("storefront")).avg(col("completedTransactions"))

Correct

The syntax shown groups the data by the storefront Column, then calculates the average value of completed sales transactions.

9. Question 9

In Azure Databricks you are about to do some ETL on a file you have received from a customer. The file contains data about people, including:

first, middle, and last names

gender

birth date

Social Security number

Salary

You discover that the file contains some duplicate records and you have been instructed to remove any duplicates. The dropDuplicates() command will more than likely create a shuffle. To help reduce the number of post-shuffle partitions which of the following commands should you run?

1 / 1 point

- spark.sql.conf.set("spark.shuffle.partitions", 8)
- spark.conf.set("spark.sql.shuffle.partitions", 8)
- spark.conf.set("spark.sql.partitions", 8)

Correct

spark.conf.set("spark.sql.shuffle.partitions", 8) is the correct syntax

10. Question 10

You need to change a column name from “dob” to “DateOfBirth” on a spark DataFrame. Which of the following syntax is valid?

1 / 1 point

.RenameColumn("dob","DateOfBirth")

.ColumnRename("dob","DateOfBirth")

.withColumnRenamed("dob","DateOfBirth")

Correct

This is correct and will rename the column “dob” to “DateOfBirth”

Test prep

Latest Submission Grade 90%

1. Question 1

A Microsoft-managed Azure Databricks workspace virtual network (VNet) exists within the customer subscription. Information exchanged between this VNet and the Microsoft-managed Azure Databricks Control Plane VNet is sent over a secure TLS connection using which Ports?

Select two from the choices below.

1 / 1 point

- [x] Port 22

Correct

The VNet and the Microsoft-managed Azure Databricks Control Plane VNet uses port 22

- [ ] Port 53
- [x] Port 5557

Correct

The VNet and the Microsoft-managed Azure Databricks Control Plane VNet uses port

- [ ] Port 443
- [ ] Port 6667

2. Question 2

Which of the following are hosted by the Azure Databricks Control Plane?

Select all that apply.

1 / 1 point

- [x] Access control lists (ACLs)

Correct

The control plane hosts security Access control lists (ACLs).

- [x] Runtime Notebooks

Correct

The control plane hosts Runtime Notebooks.

- [x] Jobs

Correct

The control plane hosts Jobs.

- [ ] Clusters

3. Question 3

In Azure Databricks using workspace access control, individual permissions determine a user's abilities. What permission must be set to allow the user the ability to change permissions?

1 / 1 point

- ( ) Edit
- (•) Manage
- ( ) Run

Correct

This permission is required to allow the user change permissions.

4. Question 4

Azure Databricks has two types of secret scopes: Key Vault-backed and Databricks-backed. These secret scopes allow you to store secrets, such as database connection strings, securely. If someone tries to output a secret to a notebook, it is replaced by which of the following?

0 / 1 point

- ( ) CONFIDENTIAL
- ( ) REDACTED
- (•) SECRET
- ( ) HIDDEN

Incorrect

Try going back and reviewing the Describe platform architecture, security, and data protection in Azure Databricks lesson.

5. Question 5

You are starting to use Azure Databricks, and you want to do specific network customizations, such as deploying Azure Databricks data plane resources in your own VNet. Which of the following will you configure?

1 / 1 point

- ( ) VNet Peering
- (•) VNet Injection
- ( ) You cannot create a custom configuration with VNets

Correct

If you're looking to do specific network customizations, you could deploy Azure Databricks data plane resources in your own VNet. In this scenario, instead of using the managed VNet, which restricts you from making changes, you "bring your own" VNet where you have full control.

6. Question 6

Which of the following features are enabled through VNet injection?

Select all that apply.

1 / 1 point

- [x] Service Endpoints

Correct

Features enabled through VNet injection include Service Endpoint.

- [ ] Managed VNet
- [x] Single-IP SNAT and Firewall-based filtering via custom routing

Correct

Features enabled through VNet injection include Single-IP SNAT and Firewall-based filtering via custom routing.

- [x] On-Premises Data Access

Correct

Features enabled through VNet injection include On-Premises Data Access.

7. Question 7

Which statement about the Azure Databricks Data Plane is true?

1 / 1 point

- (•) The Data Plane is hosted within the client subscription and is where all data is processed and stored
- ( ) The Data Plane contains the Cluster Manager and coordinates data processing jobs
- ( ) The Data Plane is hosted within a Microsoft-managed subscription

Correct

All data is processed by clusters hosted within the client Azure subscription and data is stored within Azure Blob storage and any connected Azure services within this portion of the platform architecture.

8. Question 8

In which modes does Azure Databricks provide data encryption?

1 / 1 point

- ( ) In-transit only

At-rest only

At-rest and in-transit

Correct

Data stored in Azure Storage is encrypted using server-side encryption that is seamlessly accessed by Azure Databricks. All data transmitted between the Data Plane and the Control Plane is always encrypted in-flight via TLS.

9. Question 9

What does Azure Data Lake Storage (ADLS) Passthrough enable?

1 / 1 point

Automatically mounting ADLS accounts to the workspace that are added to the managed resource group

User security groups that are added to ADLS are automatically created in the workspace as Databricks groups

Commands running on a configured cluster can read and write data in ADLS without configuring service principal credentials

Correct

Correct. Also, authentication to ADLS from Azure Databricks clusters is automatic, using the same Azure AD identity one uses to log into Azure Databricks.

10. Question 10

What is an Azure Key Vault-backed secret scope?

1 / 1 point

It is the Key Vault Access Key used to securely connect to the vault and retrieve secrets

It is a method by which you create a secure connection to Azure Key Vault from a notebook and directly access its secrets within the Spark session

Databricks secret scope that is backed by Azure Key Vault instead of Databricks

Correct

A secret scope is provided by Azure Databricks and can be backed by either Databricks or Azure Key Vault.

Test prep

Latest Submission Grade 80%

1. Question 1

Delta Lake enables you to make changes to a table schema that can be applied automatically, without the need for DDL modifications. This functionality is referred to as?

0 / 1 point

- Time Travel
- ACID Transactions
- Schema Enforcement
- Schema Evolution

Incorrect

Try going back and reviewing the Build and query a Delta Lake lesson.

2. Question 2

One of the core features of Delta Lake is performing upserts. Which of the following statements is true regarding Upsert?

1 / 1 point

- Upsert is literally TWO operations. Update / Insert
- Upsert is supported in traditional data lakes
- Upsert is a new DML statement for SQL syntax

Correct

To UPSERT means to "Update" and "Insert". In other words, UPSERT is literally TWO operations. It is not supported in traditional data lakes.

3. Question 3

What is the Databricks Delta command to display metadata?

1 / 1 point

- SHOW SCHEMA table name
- DESCRIBE DETAIL table Name
- MSCK DETAIL table name

Correct

You display metadata by using DESCRIBE DETAIL table Name.

4. Question 4

What optimization does the following command perform: OPTIMIZE Customers ZORDER BY City?

1 / 1 point

- [x] Ensures that all data backing, for example, City='London' is colocated, then rewrites the sorted data into new Parquet files.
- [ ] Ensures that all data backing, for example, City="London" is colocated, then updates a graph that routes requests to the appropriate files.
- [ ] Creates an order-based index on the City field to improve filters against that field

Correct

ZOrdering colocates related information in the same set of files.

5. Question 5

What size does OPTIMIZE compact small files to?

1 / 1 point

- [x] Around 1 GB
- [ ] Around 100 MB
- [ ] Around 500MB

Correct

The Spark optimization team determined this value to be a good compromise between speed and performance.

6. Question 6

Which of the following can be used to successfully perform an UPSERT in a Delta dataset?

1 / 1 point

- [ ] Use UPSERT INTO my-table
- [ ] Use UPSERT INTO my-table /MERGE
- [x] Use MERGE INTO my-table USING data-to-upsert

Correct

Feedback: That's the correct syntax to perform UPSERT in a Databricks.

7. Question 7

The lambda architecture is a big data processing architecture combining both batch and real-time processing methods and features an append-only immutable data source. What are features of an append-only immutable data source?

Select all that apply.

1 / 1 point

- [x] Timestamped events are appended to existing events

Correct

Timestamped events are appended to existing events (nothing is overwritten).

- [ ] Timestamped events overwrite existing events

- [x] Data is implicitly ordered by time of arrival

Correct

Data is implicitly ordered by time of arrival.

- [x] serves as system of record

Correct

Lambda features an append-only immutable data source that serves as system of record.

8. Question 8

In the Delta Lake architecture, the refinement of the data is often referred to as Bronze, Silver and Gold Tables. Which of the following tables generally contain raw data ingested from various sources?

0 / 1 point

- ( ) Silver

- (•) Gold

- ( ) Bronze

Incorrect

Try going back and reviewing the Describe Azure Databricks Delta Lake architecture lesson.

9. Question 9

What is a lambda architecture and what does it try to solve?

1 / 1 point

- ( ) An architecture that employs the latest Scala runtimes in one or more Databricks clusters to provide the most efficient data processing platform available today

- (•) An architecture that splits incoming data into two paths - a batch path and a streaming path. This architecture helps address the need to provide real-time processing in addition to slower batch computations.

- ( ) An architecture that defines a data processing pipeline whereby microservices act as compute resources for efficient large-scale data processing

Correct

The lambda architecture is a big data processing architecture that combines both batch- and real-time processing methods.

10. Question 10

What is required to specify the location of a checkpoint directory when defining a Delta Lake streaming query?

1 / 1 point

- (selected) .writeStream.format("delta").option("checkpointLocation", checkpointPath) ...
- .writeStream.format("delta").checkpoint("location", checkpointPath) ...
- .writeStream.format("parquet").option("checkpointLocation", checkpointPath) ...

Correct

Feedback: That's the correct syntax to specify the checkpoint directory on a Delta Lake streaming query.

Test prep

Latest Submission Grade 100%

1. Question 1

Stream processing is where you continuously incorporate new data into Data Lake storage and compute results. Which of the following are examples of Stream processing?

Select all that apply.

1 / 1 point

- [ ] Monthly Payroll processing
- [ ] Invoicing
- [x] Bank Card Processing

Correct

A stream of data is treated as a table to which data is continuously appended Bank Card Processing would be an example of stream processing.

- [x] IoT Device Data

Correct

A stream of data is treated as a table to which data is continuously appended IoT Device Data would be an example of stream processing.

- [x] Game play events

Correct

A stream of data is treated as a table to which data is continuously appended Game play events would be an example of stream processing.

2. Question 2

The following example creates a schema.

dataSchema = "Recorded_At timestamp, Device string, Index long, Model string, User string, _corrupt_record String, gt string, x double, y double, z double"

In SQL syntax this is referred to as which of the following?

1 / 1 point

- ( ) Data Manipulation Language(DML)
- ( ) Data Control Language (DCL)
- (•) Data Definition Language (DDL)

Correct

Here we define the schema using a DDL-formatted string (the SQL Data Definition Language).

3. Question 3

Which of the following syntax will allow you view the list of active streams in an Azure Databricks workspace?

1 / 1 point

- spark.active.streams
- spark.streams.active
- spark.active
- spark.streams

Correct

Invoking spark.streams.active will display a list of active streams.

4. Question 4

What happens if the command option ("checkpointLocation", pointer-to-checkpoint directory) is not specified?

1 / 1 point

- When the streaming job stops, all state around the streaming job is lost, and upon restart, the job must start from scratch
- It will not be possible to create more than one streaming query that uses the same streaming source since they will conflict
- The streaming job will function as expected since the checkpointLocation option does not exist

Correct

Setting the checkpointLocation is required for many sinks used in Structured Streaming. For those sinks where this setting is optional, keep in mind that when you do not set this value, you risk losing your place in the stream.

5. Question 5

When doing a write stream command, what does the outputMode("append") option do?

1 / 1 point

- The append mode replaces existing records and updates aggregates
- The append mode allows records to be updated and changed in place
- The append outputMode allows records to be added to the output sink

Correct

The outputMode "append" option informs the write stream to add only new records to the output sink. The "complete" option is to rewrite the full output - applicable to aggregations operations. Finally, the "update" option is for updating changed records in place.

6. Question 6

In Spark Structured Streaming, what method should be used to read streaming data into a DataFrame?

1 / 1 point

- spark.read
- spark.readStream
- spark.stream.read

Correct

Use the spark.readStream method to start reading data from a streaming query into a DataFrame.

7. Question 7

In Azure Databricks when creating a new user access token, the Lifetime setting of the access token can be manually set. What is the default Token Lifetime when creating a new token?

1 / 1 point

- 60 Days
- 120 Days
- Indefinite
- 30 Days
- 90 Days

Correct

The default lifetime (Days) value is 90 days.

8. Question 8

What is the purpose of “Activities” in Azure Data Factory?

1 / 1 point

- To represent a processing step in a pipeline
- To represent a data store or a compute resource that can host execution of an activity
- To link data stores or computer resources together for the movement of data between resources

Correct

Activities represent processing steps in a Data Factory pipeline.

9. Question 9

How can parameters be passed into an Azure Databricks notebook from Azure Data Factory?

1 / 1 point

- [x] Use notebook widgets
- [ ] Use the API endpoint option on a notebook
- [ ] Deploy the notebook as a web service

Correct

You can configure parameters by using widgets on the Databricks notebook. You then pass in parameters with those names via a Databricks notebook activity in Data Factory.

Test prep

Latest Submission Grade 100%

1. Question 1

What does the CD in CI/CD mean?

1 / 1 point

- Continuous Delivery
- Continuous Deployment
- Both are correct

Correct

Continuous Delivery automates your release process up to the point where human intervention is needed, by clicking a button. Continuous Deployment takes a step further by removing the human intervention and relying on automated tests to automatically determine whether the build should be deployed into production.

2. Question 2

What sort of pipeline is required in Azure DevOps for creating artifacts used in releases?

1 / 1 point

- An Artifact pipeline
- A Release pipeline
- A Build pipeline

Correct

The output of a Build pipeline is one or more artifacts that can be used within release pipelines for automated deployments.

3. Question 3

What steps are required to authorize Azure DevOps to connect to and deploy notebooks to a staging or production Azure Databricks workspace?

1 / 1 point

- Create an Azure Active Directory application, copy the application ID, then use that as the Databricks bearer token in the Databricks Notebooks Deployment step of the Release pipeline
- In the production or staging Azure Databricks workspace, enable Git integration to Azure DevOps, then link to the Azure DevOps source code repo

Create a new Access Token within the user settings in the production Azure Databricks workspace, then use the token as the Databricks bearer token in the Databricks Notebooks Deployment step of the Release pipeline

Correct

The Access Token allows you to grant access to resources within an Azure Databricks workspace without passing in user credentials.

4. Question 4

What are the two prerequisites for connecting Azure Databricks with Azure Synapse Analytics that apply to the Azure Synapse Analytics instance?

1 / 1 point

Use a correctly formatted ConnectionString and create a database master key

Add the client IP address to the firewall's allowed IP addresses list and use the correctly formatted ConnectionString

Create a database master key and configure the firewall to enable Azure services to connect

Correct

Create a database master key and configure the firewall to enable Azure services to connect.

5. Question 5

Azure Databricks is a multitenant service and in order to provide fair resource sharing to all regional customers, limits are imposed on API calls. These limits are imposed at which level?

1 / 1 point

The Subscription level

The Cluster level

The Resource group level

The Workspace level

Correct

These limits are imposed at the Workspace level.

6. Question 6

Azure Databricks is a multitenant service and to provide fair resource sharing to all regional customers, it imposes limits on API calls. What is currently the maximum number of jobs that a workspace can create in an hour?

1 / 1 point

500

- [x] 1000
- [ ] 200
- [ ] 150

Correct

The maximum number of jobs that a workspace can create in an hour is 1000.

7. Question 7

In Azure Databricks you can deploy more than one Workspace. Best practice is to use the Hub and Spoke Model. Which of the following steps should be carried out to create a best practice Hub and Spoke Model in Azure Databricks?

1 / 1 point

- [ ] Deploy each Workspace in the same VNet
- [x] Join the Workspace spokes with the central networking hub using VNet Peering

Correct

Best practice for Hub and Spoke is to join the Workspace spokes with the central networking hub using VNet Peering.

- [x] Put all the common networking resources in a central hub Vet, including the custom DNS server

Correct

Best practice for Hub and Spoke is to put all the common networking resources in a central hub Vet, including the custom DNS server.

- [ ] Join the Workspace spokes with the central networking hub using VNet Association
- [x] Deploy each Workspace in its own VNet

Correct

Best practice for Hub and Spoke is to deploy each Workspace in its own VNet.

- [ ] Put all the common networking resources in a central hub VNet but excluding the custom DNS server.

8. Question 8

Select one of the following options to make this sentence correct:

Azure Databricks guarantees by default a _____ % uptime SLA

1 / 1 point

- [ ] 99.999
- [x] 99.95

- 99.9
- 99

Correct

99.95 is the % uptime SLA that Azure Databricks guarantees by default.

9. Question 9

If mounting an Azure Data Lake Storage (ADLS) account to a workspace, what cluster feature must be used to have ACLs within ADLS applied to the user executing commands in a notebook?

1 / 1 point

- Set spark.config.adls.impersonateuser(true)
- Enable SCIM
- Enable ADLS Passthrough on a cluster.

Correct

When enabled, authentication automatically takes place in Azure Data Lake Storage (ADLS) from Azure Databricks clusters using the same Azure Active Directory (Azure AD) identity that one uses to log into Azure Databricks. Any ACLs applied at the folder or file level in ADLS are enforced based on the user's identity.

10. Question 10

What is SCIM?

1 / 1 point

- An open standard that enables users to bring their own auth key to the Databricks environment
- An open standard that enables organizations to import both groups and users from Azure Active Directory into Azure Databricks
- An optimization that removes orphaned data from a given dataset

Correct

By default, Azure Active Directory roles have no relationship with groups created inside of Azure Databricks. SCIM enables synchronizing users and groups, and synchronization is automatic after initial import.

Course practice exam

Latest Submission Grade 85%

1. Question 1

How many drivers does a Cluster have?

1 / 1 point

- [x] Only one
- [ ] Configurable between one and eight
- [ ] Two, running in parallel.

Correct

A Cluster has one and only one driver.

2. Question 2

In Azure Databricks, what type of process are the driver and the executors?

1 / 1 point

- [ ] Python processes
- [ ] C++ processes
- [x] Java processes

Correct

The driver and the executors are Java processes.

3. Question 3

How do you list files in DBFS within a notebook?

1 / 1 point

- [ ] %fs dir /my-file-path
- [ ] ls /my-file-path
- [x] %fs ls /my-file-path

Correct

Correct. You added the file system magic to the cell before executing the ls command.

4. Question 4

We can read a CVS file when using a notebook and a spark session. Which of the following can be used to view the first couple of thousand characters of a file?

1 / 1 point

- %fs dir /mnt/training/wikipedia/pageviews/
- %fs ls /mnt/training/wikipedia/pageviews/
- %fs head /mnt/training/wikipedia/pageviews/pageviews_by_second.tsv

Correct

We can use %fs head ... to view the first couple thousand characters of a file.

5. Question 5

How do you create a DataFrame object?

1 / 1 point

- Introduce a variable name and equate it to something like myDataFrameDF =
- Use the DF.create() syntax
- Use the createDataFrame() function

Correct

This approach is the correct way to create DataFrame objects.

6. Question 6

Which of the following statements describes a wide transformation?

1 / 1 point

- A wide transformation can be applied per partition/worker with no need to share or shuffle data to other workers.
- A wide transformation requires sharing data across workers. It does so by shuffling data.
- A wide transformation applies data transformation over a large number of columns.

Correct

Wide transformation shares data across workers by shuffling data between executors.

7. Question 7

Which feature of Spark of optimization is used in shuffling operations during wide transformations?

0 / 1 point

- Tungsten Record Format
- Catalyst Optimizer
- Lazy Execution

Incorrect

Try going back and reviewing Data processing in Azure Databricks.

8. Question 8

Which of the listed methods for renaming a DataFrame's column are correct?

Select two.

1 / 1 point

- [ ] df.alias("timestamp", "dateCaptured")
- [x] df.toDF("dateCaptured")

Correct

This is a valid renaming method.

- [x] df.select(col("timestamp").alias("dateCaptured"))

Correct

This is a valid renaming method.

9. Question 9

You need to change a column name from “dob” to “DateOfBirth” on a spark DataFrame. Which of the following syntax is valid?

1 / 1 point

- (•) .withColumnRenamed("dob","DateOfBirth")
- ( ) .ColumnRename("dob","DateOfBirth")
- ( ) .RenameColumn("dob","DateOfBirth")

Correct

This is correct and will rename the column “dob” to “DateOfBirth”

10. Question 10

A Microsoft-managed Azure Databricks workspace virtual network (VNet) exists within the customer subscription. Information exchanged between this VNet and the Microsoft-managed Azure Databricks Control Plane VNet is sent over a secure TLS connection using which Ports?

Select two options.

1 / 1 point

- [ ] Port 53
- [x] Port 5557

Correct

The VNet and the Microsoft-managed Azure Databricks Control Plane VNet uses port 5557.

- [ ] Port 6667
- [ ] Port 443
- [x] Port 22

Correct

The VNet and the Microsoft-managed Azure Databricks Control Plane VNet uses port 22.

11. Question 11

You are starting to use Azure Databricks and you want to do specific network customizations, such as deploying Azure Databricks data plane resources in your own VNet. Which of the following will you configure?

1 / 1 point

- (•) VNet Injection
- ( ) VNet Peering
- ( ) You cannot create a custom configuration with VNets

Correct

If you're looking to do specific network customizations, you could deploy Azure Databricks data plane resources in your own VNet. In this scenario, instead of using the managed VNet, which restricts you from making changes, you "bring your own" VNet where you have full control.

12. Question 12

In which modes does Azure Databricks provide data encryption?

0 / 1 point

- ( ) In-transit only
- ( ) At-rest and in-transit
- ( ) At-rest only

Incorrect

You didn’t select an answer.

13. Question 13

What does Azure Data Lake Storage (ADLS) Passthrough enable?

1 / 1 point

- (•) Commands running on a configured cluster can read and write data in ADLS without configuring service principal credentials.

Automatically mounting ADLS accounts to the workspace that are added to the managed resource group.

User security groups that are added to ADLS are automatically created in the workspace as Databricks groups.

Correct

Correct. Also authentication to ADLS from Azure Databricks clusters is automatic, using the same Azure AD identity one uses to log into Azure Databricks.

14. Question 14

What is the Databricks Delta command to display metadata?

1 / 1 point

DESCRIBE DETAIL tableName

SHOW SCHEMA tablename

MSCK DETAIL tablename

Correct

You display metadata by using DESCRIBE DETAIL tableName.

15. Question 15

Which of the following can be used to successfully perform an UPSERT in a Delta dataset?

1 / 1 point

Use MERGE INTO my-table USING data-to-upsert

Use UPSERT INTO my-table /MERGE

Use UPSERT INTO my-table

Correct

That's the correct syntax to perform UPSERT in a Databricks.

16. Question 16

What is a lambda architecture and what does it try to solve?

1 / 1 point

An architecture that splits incoming data into two paths - a batch path and a streaming path. This architecture helps address the need to provide real-time processing in addition to slower batch computations.

An architecture that employs the latest Scala runtimes in one or more Databricks clusters to provide the most efficient data processing platform available today

An architecture that defines a data processing pipeline whereby microservices act as compute resources for efficient large-scale data processing

Correct

The lambda architecture is a big data processing architecture that combines both batch- and real-time processing methods.

17. Question 17

What happens if the command option("checkpointLocation", pointer-to-checkpoint directory) is not specified?

1 / 1 point

When the streaming job stops, all state around the streaming job is lost, and upon restart, the job must start from scratch

The streaming job will function as expected since the checkpointLocation option does not exist

It will not be possible to create more than one streaming query that uses the same streaming source since they will conflict

Correct

Setting the checkpointLocation is required for many sinks used in Structured Streaming. For those sinks where this setting is optional, keep in mind that when you do not set this value, you risk losing your place in the stream.

18. Question 18

What's the purpose of Activities in Azure Data Factory?

0 / 1 point

To link data stores or computer resources together for the movement of data between resources

To represent a data store or a compute resource that can host execution of an activity

To represent a processing step in a pipeline

Incorrect

Try going back and reviewing Analyze streaming data and create production workloads.

19. Question 19

What sort of pipeline is required in Azure DevOps for creating artifacts used in releases?

1 / 1 point

A Build pipeline

- A Release pipeline
- An Artifact pipeline

Correct

The output of a Build pipeline is one or more artifacts that can be used within release pipelines for automated deployments.

20. Question 20

In Azure Databricks you can deploy more than one Workspace. Best practice is to use the Hub and Spoke Model. Which of the following steps should be carried out to create a best practice Hub and Spoke Model in Azure Databricks?

1 / 1 point

- [x] Join the Workspace spokes with the central networking hub using VNet Peering

Correct

Best practice for Hub and Spoke is to join the Workspace spokes with the central networking hub using VNet Peering.

- [ ] Put all the common networking resources in a central hub VNet but excluding the custom DNS server.
- [x] Deploy each Workspace in its own VNet

Correct

Best practice for Hub and Spoke is to deploy each Workspace in its own VNet.

- [x] Put all the common networking resources in a central hub Vet, including the custom DNS server

Correct

Best practice for Hub and Spoke is to put all the common networking resources in a central hub Vet, including the custom DNS server.

- [ ] Join the Workspace spokes with the central networking hub using VNet Association
- [ ] Deploy each Workspace in the same VNet