

Course Practice Exam

Latest Submission Grade 95%

1. Question 1

Which of the following is a cloud integration service that orchestrates the movement of data between various data stores?

1 / 1 point

- ☐ Azure HDInsight
- ☐ Azure Synapse Analytics
- ☒ Azure Data Factory

Correct

Azure Data Factory (ADF) is a cloud integration service that orchestrates that movement of data between various data stores. Data Engineers can create data-driven workflows in the cloud for orchestrating and automating data movement and data transformation.

2. Question 2

Azure HDInsight provides technologies for ingesting, processing, and analyzing big data to support batch processing, data warehousing, IoT, and Data Science. Which of the following technologies integrated into HDInsight provides a NoSQL database built upon Hadoop commonly used for search engines and includes automatic failover?

1 / 1 point

- ☐ Apache Storm
- ☒ Apache HBase
- ☐ Apache Kafka

Correct

Apache HBase is a NoSQL database built upon Hadoop commonly used for search engines and includes automatic failover.

3. Question 3

As part of an ETL process, data in date-formats may be modified to match the destination format. At what step in the ETL process is this carried out?

1 / 1 point

- ☒ Transformation
- ☐ Extraction
- ☐ Load

Correct

In the ETL process, data is modified during the transformation step.

4. Question 4

Which one of the following is a benefit of using the Azure SQL Database cloud-based Platform-as-a-Service (PaaS) service rather than on-premises database management systems?

1 / 1 point



Increased scalability



Increased day to day management costs



Increased functionality

Correct

PaaS increases scalability as resources can be scaled up and scaled out on demand.

5. Question 5

Which Cosmos DB API works with Graph Databases?

1 / 1 point



Gremlin



Table



MongoDB



CassandraDB

Correct

The Gremlin API enables you to perform graph queries over data.

6. Question 6

Which of the following services in Microsoft Azure are used to process data?

Select all options that apply.

0.5 / 1 point



Azure Blob Storage



Event Hub

This should not be selected

Try going back and reviewing **Services on the Microsoft Azure Data Platform**.



Stream Analytics Job

Correct

Feedback: An Azure Stream Analytics job consists of an input, query, and an output.



Synapse Analytics

7. Question 7

Which service in Azure consists of input and output pipelines, and can be used to process streaming data?

1 / 1 point



Synapse Analytics



Azure Databricks



Stream Analytics Job

Correct

An Azure Stream Analytics job consists of an input, query, and output.

8. Question 8

Data Engineers configure ingestion components of Azure Stream Analytics by configuring data inputs from sources including Azure Event Hubs, Azure IoT Hub, or Azure Blob storage.

Which of the following is a feature of Azure Event Hub?

1 / 1 point

- ☐ Partitioned consumer model
- ☐ Azure Storage
- ☒ Big data streaming service

Correct

Event hubs provide a big data streaming service of Azure. It is designed for high throughput data streaming scenarios where customers may send billions of requests per day.

9. Question 9

Nonstructured data is stored in nonrelational systems, commonly called unstructured or NoSQL systems. Which of the following are types of NoSQL databases?

Select all options that apply.

1 / 1 point

- ☒ Key-value store

Correct

A Key Value store can store key-value pairs of data in a table structure.

- ☒ Document databases.

Correct

A document database can store documents that are tagged with metadata to aid document searches.

- ☐ Relational database

- ☒ Column database

Correct

Column Databases store data based on columns rather than rows. Columns can be defined at the query's runtime, allowing flexibility in the data that's returned performantly.

- ☒ Graph database

Correct

Graph database finds relationships between data points by using a structure that's composed of vertices and edges.

10. Question 10

Which of the following is not a data job role?

1 / 1 point

- ☐ Data Analyst
- ☐ Data Administrator
- ☒ Systems Administrator

Correct

Systems administrators are responsible for maintaining and managing components such as networks, virtual machines, and other physical devices in a datacenter.

Course Practice Exam

Latest Submission Grade 95%

1. Question 1

In Azure Synapse Analytics, the Data Movement Service (DMS) coordinates and transports data between compute nodes as necessary. Azure Synapse Analytics supports which types of distributed tables.

Select all options that apply.

0.75 / 1 point



Hash

Correct

Azure Synapse Analytics supports distributed hash tables,



Replicated



Round Robin

Correct

Azure Synapse Analytics supports distributed round-robin tables.



Relational Tables

You didn't select all the correct answers

2. Question 2

Data that is continuously sent to a system for further processing in real-time is an example of which of the following data types?

1 / 1 point



Streaming data



Relational data



Batch data

Correct

Processing data in real-time, is called streaming data.,

3. Question 3

Which of the following is an example of unstructured data?

1 / 1 point



Audio Files



A Products table



A JSON File

Correct

Audio files are a type of unstructured data.

4. Question 4

Which one of the following is a benefit of using the Azure SQL Database cloud-based Platform-as-a-Service (PaaS) service rather than on-premises database management systems?

1 / 1 point

- ☐ Increased functionality
- ☒ Increased scalability
- ☐ Increased day to day management costs

Correct

PaaS increases scalability as resources can be scaled up and scaled out on demand.

5. Question 5

Which of the following is an Azure database product that works with non-relational data such as JSON data?

1 / 1 point

- ☐ Blob Storage
- ☐ SQL Server
- ☒ Cosmos DB
- ☐ Azure SQL Database

Correct

Cosmos DB is Microsoft's Azure non-relational database solution

6. Question 6

Which of the following services in Microsoft Azure are used to process data?

Select all options that apply.

0.5 / 1 point

- ☐ Event Hub
- ☒ Synapse Analytics
- ☒ Stream Analytics Job

Correct

Feedback: An Azure Stream Analytics job consists of an input, query, and an output.

- ☐ Azure Blob Storage

This should not be selected

Try going back and reviewing **Services on the Microsoft Azure Data Platform**.

7. Question 7

Which service in Azure consists of input and output pipelines, and can be used to process streaming data?

1 / 1 point

- ☒ Stream Analytics Job
- ☐ Azure Databricks
- ☐ Synapse Analytics

Correct

An Azure Stream Analytics job consists of an input, query, and output.

8. Question 8

Which of the following services can be used to ingest data into Azure Synapse Analytics?

Select all options that apply.

1 / 1 point

☐ Azure Active Directory

☒ Azure Data Factory

Correct

To ingest data, a data engineer can use either the Azure Data Factory or Azure Storage Explorer. They can also use the AzCopy Tool, PowerShell, or Visual Studio.

☒ Azure Storage Explorer

Correct

To ingest data, a data engineer can use either the Azure Data Factory or Azure Storage Explorer. They can also use the AzCopy Tool, PowerShell, or Visual Studio.

9. Question 9

Azure Data Lake Storage Gen 2 is designed to store massive amounts of data for big-data analytics. Features of Azure Data Lake include which of the following.

Select all options that apply.

1 / 1 point

☒ Hadoop compatibility

Correct

One of the features of Data Lake Storage is Hadoop compatibility.

☒ Security support for access control lists (ACLs)

Correct

One of the features of Data Lake Storage is support for access control lists.

☒ Unlimited scalability

Correct

One of the features of Data Lake Storage is unlimited scalability.

☐ Limited support for Zone Redundant storage

10.

Question 10

Which of the following is not a data job role?

1 / 1 point

☐ Data Administrator

☒ Systems Administrator

☐ Data Analyst

Correct

Systems administrators are responsible for maintaining and managing components such as networks, virtual machines, and other physical devices in a datacenter.

Course Practice Exam

Latest Submission Grade 95%

1. Question 1

Which of the following is a cloud integration service that orchestrates the movement of data between various data stores?

1 / 1 point

- ☐ Azure HDInsight
- ☐ Azure Synapse Analytics
- ☒ Azure Data Factory

Correct

Azure Data Factory (ADF) is a cloud integration service that orchestrates that movement of data between various data stores. Data Engineers can create data-driven workflows in the cloud for orchestrating and automating data movement and data transformation.

2. Question 2

Data that is continuously sent to a system for further processing in real-time is an example of which of the following data types?

1 / 1 point

- ☒ Streaming data
- ☐ Relational data
- ☐ Batch data

Correct

Processing data in real-time, is called streaming data.,

3. Question 3

Which of the following is an example of unstructured data?

1 / 1 point

- ☐ A Products table
- ☐ A JSON File
- ☒ Audio Files

Correct

Audio files are a type of unstructured data.

4. Question 4

Which of the following describes the collection of real-time telemetry data?

1 / 1 point

- ☐ Batch Processing
- ☐ Online Transaction Processing (OLTP)
- ☒ Stream Processing
- ☐ Online Analytical Processing (OLAP)

Correct

In stream processing, each new piece of data is processed when it arrives.

5. Question 5

Which of the following is an Azure database product that works with non-relational data such as JSON data?

1 / 1 point

- ☐ SQL Server
- ☒ Cosmos DB
- ☐ Azure SQL Database
- ☐ Blob Storage

Correct

Cosmos DB is Microsoft's Azure non-relational database solution

6. Question 6

Which storage solution supports role-based access control (RBAC) at the file and folder level?

0 / 1 point

- ☒ Azure Data Lake Storage
- ☒ Azure Blob storage
- ☐ Azure Disk Storage

Incorrect

Try going back and reviewing **Services on the Microsoft Azure Data Platform**.

7. Question 7

Which service in Azure consists of input and output pipelines, and can be used to process streaming data?

1 / 1 point

- ☐ Azure Databricks
- ☐ Synapse Analytics
- ☒ Stream Analytics Job

Correct

An Azure Stream Analytics job consists of an input, query, and output.

8. Question 8

Which of the following services can be used to ingest data into Azure Synapse Analytics?

Select all options that apply.

1 / 1 point

- ☒ Azure Storage Explorer

Correct

To ingest data, a data engineer can use either the Azure Data Factory or Azure Storage Explorer. They can also use the AzCopy Tool, PowerShell, or Visual Studio.

- ☒ Azure Data Factory

Correct

To ingest data, a data engineer can use either the Azure Data Factory or Azure Storage Explorer. They can also use the AzCopy Tool, PowerShell, or Visual Studio.

- ☐ Azure Active Directory

9. Question 9

Nonstructured data is stored in nonrelational systems, commonly called unstructured or NoSQL systems. Which of the following are types of NoSQL databases?

Select all options that apply.

0.8 / 1 point

- ☒ Document databases.

Correct

A document database can store documents that are tagged with metadata to aid document searches.

- ☐ Relational database

- ☒ Column database

- ☒ Graph database

Correct

Graph database finds relationships between data points by using a structure that's composed of vertices and edges.

- ☒ Key-value store

Correct

A Key Value store can store key-value pairs of data in a table structure.

You didn't select all the correct answers

10. Question 10

Which of the following is not a data job role?

1 / 1 point

- ☐ Data Administrator

- ☐ Data Analyst

- ☒ Systems Administrator

Correct

Systems administrators are responsible for maintaining and managing components such as networks, virtual machines, and other physical devices in a datacenter.

Course practice exam

Latest Submission Grade 93.89%

1. Question 1

Which of the following are benefits of Azure cloud-based data storage solutions?

Select all options that apply.

1 / 1 point



Multiple Data Types

Correct

Azure can store almost any type of data you need. It can handle video files, text files, and even large binary files like virtual hard disks. It also has many options for your relational and NoSQL data.



Increased Capital Expenditure



Automated backup and Recovery

Correct

Automated backup and recovery mitigate the risk of losing your data if there is any unforeseen failure or interruption.



Global Replication

Correct

Replication across the globe copies your data to protect it against any planned or unplanned events, such as scheduled maintenance or hardware failures. You can choose to replicate your data at multiple locations across the globe.

2. Question 2

Which of the following statements is true of Resource Groups, Storage Accounts, and Azure Cosmos DB?

1 / 1 point



Azure Cosmos DB can be included in an Azure Resource Group but cannot be included in an Azure Storage Account.



Azure Cosmos DB can be included in an Azure Storage Account in an Azure Resource Group

Correct

A storage account is an Azure resource and is included in a resource group. Data services like Azure SQL and Azure Cosmos DB are managed as independent Azure resources and can be included in an Azure Resource Group but cannot be included in a storage account.

3. Question 3

Transactions are often defined by a set of four requirements, referred to as ACID guarantees.

What does ACID stand for?

Select all options that apply.

1 / 1 point



Consistency

Correct

Consistency ensures that the data is consistent both before and after the transaction.



Isolation

Correct

Isolation ensures that one transaction is not impacted by another transaction.



Automatic



Dynamic



Concurrent



Atomicity

Correct

Atomicity means a transaction must execute exactly once and must be atomic; either all of the work is done, or none of it is. Operations within a transaction usually share a common intent and are interdependent.



Durability

Correct

Durability means that the changes made due to the transaction are permanently saved in the system. Committed data is saved by the system so that even in the event of a failure and system restart, the data is available in its correct state.

4. Question 4

Let's say you have two video files stored as blobs. One of the videos is business-critical and requires a replication policy that creates multiple copies across geographically diverse datacenters. The other video is non-critical, and a local replication policy is sufficient. Which of the following options satisfies both data diversity and cost sensitivity consideration?

1 / 1 point



Create a single storage account that makes use of local-redundant storage (LRS) and host both videos from here.



Create a single storage account that makes use of geo-redundant storage (GRS) and host both videos from here.



Create two storage accounts. The first account makes use of Geo-redundant storage (GRS) and hosts the business-critical video content. The second account makes use of local-redundant storage (LRS) and hosts the non-critical video content.

Correct

Try going back and reviewing **Store Data in Microsoft Azure**.

5. Question 5

Storage for VHDs is used to provide durable disks for Azure Virtual Machines that utilize which type of Blob Storage?

1 / 1 point



Block Blobs



Page Blob

Feedback. Page blobs are used primarily as the backing storage for the VHDs to provide durable disks for Azure Virtual Machines.



Append Blobs

Correct

6. Question 6

What is the maximum size of a block blob storing text or binary files?

1 / 1 point



5 TB



500 GB



5 GB

Correct

Block blobs are used to hold text or binary files up to ~5 TB (50,000 blocks of 100 MB) in size.

7. Question 7

Which of the following are features of the Azure Queue service?

Select all options that apply.

0.75 / 1 point



Queue messages can be up to 64 kilobytes in size, and a queue can contain millions of messages.

Correct

Queue messages can be up to 64 kilobytes in size, and a queue can contain millions of messages.



Queues are used to store lists of messages to be processed synchronously.



Queues are used to store lists of messages to be processed asynchronously.

Correct

Queues are used to store lists of messages to be processed asynchronously.



It can store and retrieve messages.

You didn't select all the correct answers

8. Question 8

You have added the required client libraries to your application and are ready to connect to your Azure storage account. To work with data in a storage account, your app will need two pieces of data. What are these pieces of data?

1 / 1 point



Azure Active Directory Account



REST API endpoint

Correct

To work with data in a storage account, your app will need a REST API endpoint.



Access key

Correct

To work with data in a storage account, your app will need an Access key.

9. Question 9

Azure Storage provides a REST API to work with the containers and data stored in each account. To work with data in a storage account, what pieces of data does your app need?

Select all options that apply.

1 / 1 point



Private access key



Access key

Correct

Azure Storage provides a REST API to work with the containers and data stored in each account. To work with data in a storage account, your app will need two pieces of data an access key and a REST API endpoint



REST API Endpoint

Correct

Azure Storage provides a REST API to work with the containers and data stored in each account. To work with data in a storage account, your app will need two pieces of data an access key and a REST API endpoint

- ☐ Public access key

10. Question 10

Microsoft Azure Storage is a managed service that provides durable, secure, and scalable storage in the cloud. Azure Files enables you to set up highly available network file shares that can be accessed using the standard Server Message Block (SMB) protocol.

Which of the following are common scenarios where file shares can be used?

Select all options that apply.

0.3333333333333333 / 1 point

- ☐ Sharing data between on-premises applications and Azure VMs to allow migration of apps to the cloud over time.

- ☒ Sharing data between on-premises and Azure VM's to allow the migration of apps to the cloud instantly

This should not be selected

Try going back and reviewing **Connect an app to Azure Storage**.

- ☒ Storing shared configuration files for VMs, tools, or utilities so that everyone is using the same version.

Correct

Multiple VMs can share the same files with both read and write access.

11. Question 11

When connecting an app to multiple storage accounts, what will your app require?

1 / 1 point

- ☒ An Access key for each storage account

- ☐ A Shared Access Key

- ☐ A single Access key to gain access to all storage accounts

Correct

Each storage account has two unique access keys that are used to secure it. If your app needs to connect to multiple storage accounts, your app will require an access key for each storage account.

12. Question 12

By default, storage accounts accept connections from clients on any network. To limit access to selected networks you must first change this default action.

Which of the following can you use to restrict access to selected networks?

1 / 1 point

☒ Ranges of Ip addresses

Correct

You can restrict access to a range of IP Addresses.

☒ Specific IP Addresses

Correct

You can restrict access to specific IP Addresses.

☐ Resource Group

☒ Virtual Networks

Correct

You can restrict access to a Virtual Network.

13. Question 13

Azure Data Lake Storage Gen2 is built on which of the following storage services?

1 / 1 point

☒ Azure Blob Storage

☐ Azure Tables

☐ Azure File Storage

☐ Azure Disks

Correct

Azure Data Lake Storage Gen2 is built on Azure Blob Storage.

14. Question 14

The organization you work at has two types of data:

1. Private and proprietary

2. For public consumption.

When considering Azure storage accounts, which of the following approaches is best to meet your data diversity requirement?

1 / 1 point

☒ Enable virtual directories for the proprietary data and not for the public data with separate storage accounts for the proprietary and public data.

- ☐ Locate the organization's data in a datacenter in the required country or region with one storage account for each region.
- ☐ Locate the data in a datacenter with the strictest data regulations to ensure that regulatory requirements thresholds have been met. In this way, you need only one storage account for managing all data. This will reduce data storage costs.

Correct

You need one storage account for every group of settings that you want to apply to your data.

15. Question 15

Which of the following statements is true of storage accounts and blob containers?

1 / 1 point

- ☐ Every blob lives inside a blob container. You can only store a limited number of blobs in a container and a limited number of containers in a storage account.

☒ Every blob lives inside a blob container. You can store an unlimited number of blobs in a container and an unlimited number of containers in a storage account.

- ☐ Every blob lives inside a blob container. You can store a limited number of blobs in a container and an unlimited number of containers in a storage account.

Correct

Every blob lives inside a blob container. You can store an unlimited number of blobs in a container and an unlimited number of containers in a storage account.

Course practice exam

Latest Submission Grade 83.33%

1. Question 1

Much of the functionality of Azure Data Factory appears in Azure Synapse Analytics as a feature called pipelines. You can use it to integrate data pipelines between which of the following?

1 / 1 point

☒ SQL Pools

Correct

Much of the functionality of Azure Data Factory appears in Azure Synapse Analytics as a feature referred to as Pipelines, which enables you to integrate data pipelines between SQL Pools, Spark Pools and SQL Serverless.

☐ Apache Hive

☒ SQL Serverless

Correct

Much of the functionality of Azure Data Factory appears in Azure Synapse Analytics as a feature referred to as Pipelines, which enables you to integrate data pipelines between SQL Pools, Spark Pools and SQL Serverless.

☒ Spark Pools

Correct

Much of the functionality of Azure Data Factory appears in Azure Synapse Analytics as a feature referred to as Pipelines, which enables you to integrate data pipelines between SQL Pools, Spark Pools and SQL Serverless.

2. Question 2

Microsoft Azure provides a variety of data platform services that enables you to perform different types of analytics. Predictive analytics can be implemented through which of the following features?

Select all options that apply.

0.75 / 1 point

☒ Machine Learning Services

Correct

Microsoft Azure provides a variety of data platform services that enables you to perform different types of analytics including predictive analytics within HDInsight, Azure Databricks or Machine Learning Services.

☒ Azure Databricks

Correct

Microsoft Azure provides a variety of data platform services that enables you to perform different types of analytics including predictive analytics within HDInsight, Azure Databricks or Machine Learning Services.

☐ Azure Data Lake Storage Gen2

☐ HDInsight

You didn't select all the correct answers

3. Question 3

When setting up Azure Data factory through the Azure Portal. Which of the following pieces of Information are required?

Select all options that apply.

1 / 1 point

☒ Region

Correct

When setting up Azure Data Factory you are required to supply the Region Information.

☒ Subscription

Correct

When setting up Azure Data Factory you are required to supply the subscription in which the ADF instance is created.

☒ Resource Group name

Correct

When setting up Azure Data Factory you are required to supply the Resource Group Name.

☐ Management Group Name

☒ Instance Name

Correct

When setting up Azure Data Factory you are required to supply the name of the Azure Data Factory instance.

4. Question 4

Which of the following is true in respect of Azure data factory datasets?

0 / 1 point

☐ A dataset is a named view that points or references the data.

☒ A dataset is a table that holds a copy of the data for the destination.

- ☐ A dataset does not exist in Azure Data factory
- ☐ A dataset is a table that holds a copy of the data from the source.

Incorrect

Try going back and reviewing **Set up and manage Azure Data Factory**.

5. Question 5

Data Factory offers three types of Integration Runtime.

- Azure
- Self-hosted
- Azure-SSIS

Which of these provides support for both Public Networks and Private Networks?

Select all options that apply.

1 / 1 point

- ☐ Azure only
- ☐ Azure and Self Hosted

☒ Self Hosted and Azure-SSIS

- ☐ Azure and Azure-SSIS
- ☐ Azure-SSIS only
- ☐ Self Hosted only

Correct

Self Hosted and Azure-SSIS support both Public and Private networks.

6. Question 6

Which Azure Data Factory component contains the transformation logic or the analysis commands of the Azure Data Factory's work?

1 / 1 point

- ☐ Linked Services

☒ Activities

- ☐ Pipelines
- ☐ Datasets

Correct

Activities contains the transformation logic or the analysis commands of the Azure Data Factory's work.

7. Question 7

Azure Data Factory can call on compute resources to process data by a specific data platform service that may be better suited to the job. Which compute resource can carry out the following activities, Spark, MapReduce and Hadoop Streaming activities?

1 / 1 point

- ☐ Azure Databricks
- ☐ Azure Machine Learning
- ☐ Azure Batch

☒ On-demand HDInsight cluster

Correct

On-demand HDInsight cluster or your own HDInsight cluster can carry out activities such as Hive, Pig, Spark, MapReduce, Hadoop Streaming.

8. Question 8

You want to ingest data from a SQL Server database hosted on an on-premises Windows Server. What integration runtime is required for Azure Data Factory to ingest data from the on-premises server?

0.6666666666666666 / 1 point

- ☐ Azure Integration Runtime

☒ Azure-SSIS Integration Runtime

This should not be selected

Try going back and reviewing **Petabyte-scale ingestion with Azure Data Factory or Azure Synapse Pipeline**.

☒ Self-Hosted Integration Runtime

Correct

A self-hosted integration runtime can run copy activities between a cloud data store and a data store in a private network. It also can dispatch transform activities against compute resources in an on-premises network or an Azure virtual network.

9. Question 9

To create Data Factory instances, the user account that you use to sign into Azure must be a member of which of the following roles?

Select all options that apply.

0.75 / 1 point

☒ The Owner role

Correct

The Owner role grants full access to manage all resources, including the ability to assign roles in Azure RBA.

☐ Managed Application Operator Role

☒ The Contributor Role

Correct

The Contributor Role grants full access to manage all resources, but does not allow you to assign roles in Azure RBAC, manage assignments in Azure Blueprints.

☒ Monitoring Contributor

This should not be selected

Try going back and reviewing **Petabyte-scale ingestion with Azure Data Factory or Azure Synapse Pipeline**.

10. Question 10

Azure Data Factory provides a range of methods to perform transformations. You need to perform transformations without writing code. Which transformations would you use in this scenario?

0 / 1 point

- ☐ Stored procedure activity
- ☐ HDInsight Pig activity
- ☐ Mapping Data Flows

☒ Azure Machine Learning Studio

Incorrect

Feedback: In this scenario, you would not use Azure Machine Learning Studio. Try going back and reviewing **Perform code-free transformation at scale with Azure Data Factory or Azure Synapse Pipeline**.

11. Question 11

Which transformation is used to load data into a data store or compute resource?

1 / 1 point

☐ Source

☒ Sink.

☐ Window

☐ Surrogate key

Correct

A Sink transformation allows you to choose a dataset definition for the destination output data. You can have as many sink transformations as your data flow requires.

12. Question 12

You want to load a dimension table in Azure Synapse from source data in your Azure Synapse database using mapping data flows. Which of the following options would you choose in Synapse Studio Hub to create a Type 1 SCD in a Mapping Data Flow pipeline activity?

1 / 1 point

☒ Develop

☐ Integrate

☐ Manage

☐ Data

Correct

The Develop hub contains your SQL scripts, notebooks, data flows, and Spark job definitions.

13. Question 13

In Azure Data factory there are many activities that are possible in a pipeline. These activities are grouped into 3 categories. Which of the following are categories of activities?

Select all options that apply.

0.5 / 1 point

☐ Data transformation activities

☒ Control activities

Correct

Examples of control flow activities are 'get metadata', 'ForEach', and 'Execute Pipeline'.

☒ Branching Activities

This should not be selected

Try going back and reviewing **Orchestrate data movement and transformation in Azure Data Factory or Azure Synapse Pipeline**.

☒ Data movement activities

Correct

The Copy Activity in Data Factory copies data from a source data store to a sink data store.

14. Question 14

In Azure Data Factory to get a view of the history of debug runs, or see all the active debug runs, you can navigate to the monitor tab. Azure Data Factory service keeps Debug run history for how many days?

1 / 1 point



15 days



30 Days



365 Days



5 Days

Correct

Azure Data Factory service keeps Debug run history for 15 Days.

15. Question 15

To lift and shift an existing SSIS workload, you can create an Azure-SSIS IR to natively execute SSIS packages. The location for your Azure-SSIS IR is important to achieve high performance in your extract-transform-load (ETL) workflows. Select which of the following statements are true.

1 / 1 point



The location of your Azure-SSIS IR must be the same as the location of your data factory.



The location of your Azure-SSIS IR does not need to be the same as the location of your data factory.

Correct

To lift and shift an existing SSIS workload, The location of your Azure-SSIS IR does not need to be the same as the location of your data factory.

16. Question 16

When you migrate your database workloads from SQL Server on premises to Azure SQL database services, you may have to migrate SSIS packages as well. The first step required is to perform an assessment of your current SSIS packages to make sure that they are compatible in Azure. Which of the following tools would be most appropriate to check for compatibility issues?

1 / 1 point



SQL Server Data Tools (SSDT)



Data Migration Assistant (DMA)



SQL Server Management Studio (SSMS)

Correct

The first step required is to perform an assessment of your current SSIS packages to make sure that they are compatible in Azure. You can use the Data Migration Assistant (DMA) to perform an assessment of the SSIS packages that exist and identify any compatibility issues with them.

17. Question 17

An SSIS Project is the unit of deployment for SSIS solutions. In which version of SQL Server was SSIS Projects introduced?

1 / 1 point

- ☐ SQL Server 2008.
- ☐ SQL Server 2016

☒ SQL Server 2012

Correct

SSIS Projects were introduced in SQL Server 2012 and is the unit of deployment for SSIS solutions.

18. Question 18

Azure Data Factory allows you to configure a Git repository with either Azure Repos or GitHub. Which of the following are advantages of using GIT Integration?

Select all options that apply.

1 / 1 point

☒ Better CI/CD

Correct

GIT integration makes certain actions easier such as Configuring release pipeline to trigger automatically as soon as there are any changes made to your 'dev' factory.

☐ PowerShell or an SDK changes are published directly to GIT.

☒ Source control

Correct

Using GIT Source Control provides benefits such as the ability to track/audit changes and revert changes that introduced bugs.

☒ Partial saves

Correct

Configuring a git repository allows you to save changes, letting you only publish when you have tested your changes to your satisfaction.

19. Question 19

Mapping data flows allow you to build code-free data transformation logic that runs at scale. When building your logic, you can turn on a debug session to interactively work with your data. Turning on Debug creates which of the following?

1 / 1 point

- ☐ MS SQL Cluster
- ☐ Data Lake Storage Account
- ☒ Spark Cluster
- ☐ Blob Storage

Correct

When building your logic, you can turn on a debug session to interactively work with your data using a live Spark cluster.

20. Question 20

In Azure Data factory what is a supported connector for built-in parameterization?

1 / 1 point

- ☐ Azure Data Lake Storage Gen2.

☒ Azure Synapse Analytics

- ☐ Azure Key Vault

Correct

Azure Synapse Analytics is a supported connector for built-in parameterization for Linked Services in Azure Data Factory.